

Kontablo: A Graph-Based Universal Accounting Ontology for the M2M Agentic Economy*

Christian Luciani[†]

Version 1.8 — June 2026

Abstract

Kontablo¹ is an open, graph-based universal accounting ontology anchored to International Financial Reporting Standards (IFRS) and mapped across all 195 sovereign jurisdictions through a universal IFRS-anchored layer, with statutory chart-of-accounts overlays for the 60 jurisdictions that mandate a national chart, providing a deterministic, jurisdiction-agnostic subledger for both multinational ERP consolidation and the Machine-to-Machine (M2M) Agentic Economy.

The global accounting ecosystem remains fragmented across incompatible national standards — an “Accounting Babel” that consumes an estimated 30–40% of corporate finance teams’ time in reconciliation and sustains a multi-billion-dollar financial-close software market, while leaving high-frequency autonomous transactions without a common execution layer. The status quo is not the absence of automation but its fragmentation: bespoke, point-to-point integrations between every pair of systems rather than one shared semantic standard. As AI agents begin to execute financial operations via Agent Payments (AP2) and Agent2Agent (A2A) protocols, they need a semantic execution layer that reporting standards (XBRL) and payment messaging (ISO 20022) do not provide.

Kontablo addresses this gap through three core contributions: (1) a Level 3 minimum-core taxonomy of 30 universal accounts covering an estimated 94% of routine transaction volume by posting count (~99% with a 34-account extended core), reproducible from a committed benchmark rather than asserted; (2) a **Three-Tier Resolution Strategy** combining deterministic exact-code lookups and keyword rules with a confidence-scored semantic AI fallback, bounded by a Deterministic

*Cite as: Luciani, C. (2026). *Kontablo: A Graph-Based Universal Accounting Ontology for the M2M Agentic Economy*. Preprint v1.8. github.com/ChristianLuciani/accounting-esperanto. BibTeX: `@misc{luciani2026kontablo, author={Luciani, Christian}, title={Kontablo: A Graph-Based Universal Accounting Ontology for the {M2M} Agentic Economy}, year={2026}, howpublished={Preprint v1.8}}`

[†]Chief Architect, Kontablo Project. Independent Researcher, Cuenca, Ecuador. ORCID: 0000-0002-6955-5384. Contact: cluciani@gmail.com.

¹**Etymology.** *Kontablo* is a neologism coined for this project. Its roots are Esperanto: *konto* (a financial account) and the cognate *kontado* (accountancy), both tied to *kalkuli* (to calculate, to reckon) (Reta Vortaro, <https://vortaro.net>). The form deliberately echoes the Romance *contable/contabile* (“accountant”, “of accounts”), making the term immediately recognizable to Spanish-, Portuguese-, Italian-, and French-speaking readers while remaining neutral in spirit. We read *Kontablo* as “a system of accounts”: a sustained, calculable accounting discipline rather than a single ledger — fitting for a universal standard built for a hybrid economy of human and AI agents. (Morphology note: regular Esperanto accountancy derivations use *-ado/-eco*; *Kontablo* is a deliberate coinage, not a standard inflection.)

Boundary Library that eliminates the classical accounting hallucination class at the harness level; and (3) a **Co-responsibility Architecture (CRA)** pairing every AI mapping proposal with a mandatory human review pathway and an append-only audit trail, keeping legal accountability with the human operator. On a synthetic validation matrix built from the ontology’s own committed code sets — 75 entities, 68 jurisdictions, fourteen languages across five writing systems, four IAS 29 dual-rate hyperinflation stress tests — the deterministic tiers resolve 97% of entries, out-of-coverage instruments escalate to human review rather than mismapping, and the CRA quarantines a mechanically wired injected-error catalog. Tier-1 coverage is exercised against primary-source-cited statutory charts in 56 of the 60 statutory-chart jurisdictions (full enumeration and construction protocol in the Evaluation section); the run also closed a corrective loop on the standard, surfacing four latent code collisions later fixed against primary sources. Kontablo is published under the Business Source License 1.1, converting to Apache 2.0 after four years.

Keywords: Accounting Standards · Knowledge Graphs · IFRS · Agentic Economy · M2M Transactions · ERP Interoperability · Co-responsibility Architecture · AI Governance · Human-in-the-Loop · Algorithmic Accountability

1 Introduction

1.1 The Three Structural Crises of Modern Accounting

Accounting systems are often viewed merely as recording tools, but they represent the actual *language* of capital. When this language is fragmented, capital becomes inefficient. Current multinational accounting environments suffer from three core structural crises that legacy architectures cannot resolve:

1. **The M2M Agentic Void:** We are entering the era of the “Agentic Economy,” where autonomous AI agents (negotiating via frameworks like the Agent2Agent (A2A) protocol and the Agent Payments Protocol (AP2), both now governed by standards bodies) execute financial transactions at high frequencies. The IMF has identified the core tension as “modern payment systems are built on deterministic logic, while agentic AI operates probabilistically” (Davidovic & Tourpe, IMF Notes 2026/004). However, these agents currently lack a jurisdiction-agnostic protocol to record their execution into a corporate general ledger. An AI agent cannot halt mid-transaction to parse localized regulatory variances, such as a Vietnamese VAS rulebook or an Arabic SOCPA tax provision; it requires a singular, deterministic ID to book its actions. Without a standardized subledger layer, the transaction velocity enabled by AI agents will create an unmanageable reconciliation bottleneck.
2. **Cross-Border ERP Fragmentation (*The Accounting Babel*):** Financial consolidation remains heavily manual and fragmented, and enterprises spend billions annually to contain the problem. As a conservative proxy for its scale, the global financial close software market was valued at roughly **USD 8.8 billion in 2025** (projected USD 18.5 billion by 2034), and the narrower financial consolidation software segment at approximately **USD 3.2 billion** (Business Research Industry, 2026; Market Research, 2025) — expenditure that exists primarily to mitigate the reconciliation burden described here, and which excludes the far larger internal

labor cost. At the data-quality level, Gartner estimates poor data quality costs the average organization **USD 12.9 million per year** (Gartner, 2021). Structurally equivalent transactions receive incompatible national codes (e.g., Mexico’s SAT codes vs. France’s PCG) and contrasting tax nomenclatures. This fragmentation prevents global enterprises from achieving a unified, real-time balance sheet. According to industry benchmarks (Gartner, 2024; Deloitte, 2025), corporate finance teams devote between **30% and 40% of their operational time** to repetitive manual reconciliation, transaction matching, and data cleansing. This manual intervention results in a persistent entry **error rate of 3% to 25%** (Fintech Global, 2025), creating significant audit risks and financial leakage.

3. **Hyperinflationary and Dual-Rate Divergence:** In volatile economies like Venezuela (VES) or Argentina (ARS), traditional ERP constraints statically tethered to a single exchange rate fail to represent economic reality. Companies must operate against official government rates and parallel market rates simultaneously, with no native ledger protocol to reconcile these diverging financial narratives at the consolidation level. This requires dynamic, parallel-rate tracking and restatements that current relational databases cannot represent cleanly.

1.2 The IFRS Implementation Gap

The International Financial Reporting Standards (IFRS) have achieved near-universal conceptual adoption: the IFRS Foundation publishes jurisdiction profiles for 169 jurisdictions where IFRS is required or permitted for public companies (IFRS Foundation, 2025). Yet, IFRS lacks a **computational execution layer**. While a CPA in Brazil and one in Japan agree on the conceptual definition of a “Trade Receivable,” their respective localized ERP systems implement that concept using completely different numerical schemas, database structures, and tax metadata tags. IFRS regulates reporting disclosures, not the execution of the transactional general ledger itself.

1.3 Research Contribution

Terminology: “ontology” in this paper. Throughout this document, **ontology** is used in its information-science sense, not its philosophical one. We do not mean the metaphysical study of being. We mean a *formal, machine-readable specification of the concepts in a domain and the relationships between them* — the established usage in knowledge representation and the semantic web (Gruber, 1993; Spivak & Kent, 2012). A Kontablo account is an ontology element in exactly this sense: a uniquely identified concept (a UUID) carrying typed attributes (nature, statement, jurisdiction overlay) and explicit, queryable relationships to other concepts. The artifact is a computational object an agent can reason over, not a philosophical claim about the nature of economic reality.

Kontablo bridges this gap by creating an open, graph-based universal accounting ontology. Unlike prior work on XBRL—which is a disclosure and reporting taxonomy, not a transaction-execution layer (we treat the two as complementary; Section 2)—Kontablo defines a precise “Level 3” taxonomy of 30 universal accounts. This paper makes the following contributions:

1. A formal graph-based data model for IFRS-anchored universal accounts, resolving tree-hierarchy limitations.

2. A three-tier mapping resolution strategy that maps localized ERP trees into the universal graph.
3. A multilingual semantic mapping engine with a Co-responsibility Architecture (CRA) for auditing human overrides.
4. A mass-consolidation simulation engine proving coverage across 195 sovereign jurisdictions, including a dual-currency hyperinflation stress test.
5. A roadmap for two-way API integration with the top 10 global commercial ERP systems.

The three resolution tiers, up front. Because the rest of this paper refers constantly to “Tier 1/2/3,” we define them here; they are developed fully in Section 8. Each tier is an attempt to resolve one local account (a code and a name in some jurisdiction) to a universal Level 3 node, and the engine descends the tiers only on failure:

- **Tier 1 — Exact code lookup** (deterministic, confidence 1.0): the local statutory code matches a committed entry for that jurisdiction. *Example:* Mexican SAT code 101 → `asset.current.cash`.
- **Tier 2 — Deterministic keyword rules** (deterministic, confidence ~ 0.8 – 0.9): no code match, but an auditable multilingual rule fires on the account name. *Example:* a Vietnamese line named “Tien mat” (cash) → `asset.current.cash`.
- **Tier 3 — Semantic AI fallback** (stochastic, confidence 0.5–0.85): neither deterministic tier resolves, so an LLM proposes a node *from the fixed set of existing universal nodes only* — the Deterministic Boundary Library rejects any proposal outside the graph and routes it to human review. *Example:* “Provisión para Contin-gencias Laborales” with no matching code or keyword → a proposed node plus a confidence score, gated before posting.

The decisive property is that Tiers 1 and 2 are deterministic and verifiable, Tier 3 is the only stochastic step, and *no* tier can invent a node that does not already exist in the ontology.

Organization of this paper. Sections 2 and 5 establish the historical precedents and the graph-versus-tree data model; Section 6 gives the formal (category- and sheaf-theoretic) account of the translation and its reconciliation obstruction. Section 8 specifies the three-tier resolution strategy and the Deterministic Boundary Library, and Section 9 frames the system as a model-plus-harness and locates where residual error can live. Sections 10 and 12 cover the agent-native interfaces and the Co-responsibility Architecture (including its scope and enforcement boundary, Section 12.4). Section 13 reports the reproducible validation; the labor-market and conclusion sections close the paper. Readers interested only in the engineering may read Sections 5, 8, 9, and 13 as a self-contained unit.

2 The Historical Precedents of Digital Accounting

2.1 A History of Ledger Standardization

The fragmentation of accounting systems is not a modern software glitch; it is the legacy of national economic organization. The earliest formal attempts to standardize ledger structures date back to the mid-20th century, notably the German *Gemeinschaftsrahmen* of 1937 and the French *Plan Comptable Général* (PCG) of 1947. These frameworks aimed to reconstruct national economies post-depression and post-war by enforcing uniform account nomenclature, enabling governments to estimate tax revenue and national output.

During the late 20th century, this model spread globally. Spain adopted the *Plan General de Contabilidad* (PGC), which subsequently influenced the national charts of accounts (known as *Planes Únicos de Cuentas*, or PUCs) across Latin America, including Colombia, Peru, and Chile. Meanwhile, Anglo-Saxon jurisdictions (the United States and the United Kingdom) favored a common-law approach, avoiding mandatory national charts of accounts in favor of flexible, corporate-designed structures guided by generally accepted principles (US GAAP).

2.2 Why Accounting Babel Has Persisted for Decades

Despite the obvious efficiency gains of a single accounting language, the “Accounting Babel” problem has persisted for over eighty years due to three structural barriers:

1. **National Sovereignty and Fiscal Policy Enforcement:** A chart of accounts is not merely a reporting ledger; it is a regulatory control loop. Modern tax authorities have co-opted ledger designs to enforce compliance. For example, Mexico’s SAT (Servicio de Administración Tributaria) requires companies to tag every transaction with SAT classification codes to match electronic invoicing metadata (*Comprobante Fiscal Digital por Internet*, or CFDI). Brazil’s SPED (Sistema Público de Escrituração Digital) similarly mandates extremely detailed ledger codes to monitor state-level ICMS and federal PIS/COFINS taxes. Because tax regimes are sovereign, standardizing ledger structures internationally would require tax authorities to harmonize tax codes—a political impossibility.
2. **Legacy ERP Vendor Economics and Customization Lock-in:** The global enterprise software market is dominated by legacy ERP systems (SAP S/4HANA, Oracle NetSuite, Microsoft Dynamics) whose business models thrive on customization complexity. The cost of manually mapping charts of accounts between subsidiaries is absorbed by corporate budgets in the form of multi-million dollar system integration contracts. Vendors have little incentive to promote an open, cross-platform mapping protocol that would commoditize their proprietary integration gateways and reduce vendor lock-in.
3. **The Paper-Compliance Paradigm:** Accounting standards like IFRS and US GAAP were designed under the assumption of periodic, human-conducted audits of printed financial statements. Consequently, standard-setters focused on *disclosure templates* (reporting) rather than *transaction-level data structures* (execution). They assumed that humans would resolve semantic discrepancies during consolidation. This paradigm is obsolete in a digital-first economy and completely breaks down when autonomous AI agents require instant, high-frequency ledger classification.

2.3 The Reporting vs. Execution Gap

Prior attempts to solve this fragmentation have failed because they target the wrong layer of the financial stack:

- XBRL (eXtensible Business Reporting Language):** Introduced in the early 2000s, XBRL revolutionized financial disclosure by creating machine-readable tags for corporate reports. However, XBRL is strictly a *post-hoc reporting* standard. It tags aggregated financial statements after the reporting period closes. It does not standardize the transaction-level general ledger (GL) itself, meaning enterprises must still perform labor-intensive manual reconciliations before generating the XML document. We stress that this is a difference of *layer*, not a defect, and that Kontablo is **complementary to XBRL rather than a replacement for it**: because every Kontablo node carries an IFRS taxonomy tag, a Kontablo-consolidated ledger maps directly onto the IFRS/XBRL taxonomy elements that disclosure regimes such as ESEF (the EU’s mandated inline-XBRL format) consume. Kontablo produces the clean, IFRS-tagged, transaction-level data upstream; XBRL/ESEF emits the disclosure document downstream. The two compose — Kontablo node → IFRS tag → XBRL element → iXBRL filing — so adopting Kontablo does not displace an existing XBRL reporting pipeline; it removes the manual reconciliation that currently sits between the GL and the filing.
- EDIFACT and ISO 20022:** These standards represent the peak of financial messaging for payments. ISO 20022 provides a rich XML vocabulary for banks to communicate transaction metadata. Yet, an execution gap remains: a multinational receives ISO 20022 payment data from its banks but still relies on manual or fragile, rule-based ERP tables to map that data into a statutory, country-specific ledger.




Figure 1: Historical evolution of financial standards. While communication (payment messaging) and disclosure (XBRL) standardized over time, the execution layer — general-ledger chart-of-accounts mappings — remained fragmented at the national tree level until Kontablo’s graph ontology.

Kontablo identifies that the failure of previous standards to achieve “Real-Time Universal Execution” is rooted in this lack of a computational bridge between the local statutory tree and the global analytical graph, leaving a multi-billion dollar reconciliation gap.

2.4 The Incumbent Answer: ERP Homogenization — and the Estate It Leaves Behind

It is fair to ask how the largest multinationals already cope, given that they consolidate thousands of entities daily. The answer is instructive: they solve the Babel problem by **eliminating it through homogenization** rather than translation. Large groups standardize every subsidiary onto a single ERP instance with a single group chart of accounts and parallel statutory/management ledgers. SAP S/4HANA’s Universal Journal stores all postings in one unified table (ACDOCA), enabling multi-GAAP and multi-currency reporting within one instance; Oracle’s ledger architecture maintains primary (e.g. US-GAAP) and secondary (e.g. IFRS) ledgers in parallel within a single instance, with automated conversion journals between them (SAP, 2024; Oracle, 2024). This works — for those who can afford it. It presupposes a multi-year, multi-million-dollar program to force every acquired or regional entity onto the parent’s instance and chart, and it is precisely the customization lock-in described above: the group’s semantic model now lives inside one vendor’s proprietary schema.

Kontablo targets the estate homogenization leaves behind — which is the common case, not the exception. M&A-grown groups operating mixed ERPs, SMEs that cannot fund a single-instance rollout, and any enterprise unwilling to surrender its semantics to a vendor schema all face heterogeneous charts that must *interoperate* rather than be replaced. Two consequences follow directly from the tree-to-graph protocol of Section 5.3:

- **Vendor-lock-in reduction.** When an enterprise’s account semantics are expressed as edges to an *open* universal node set rather than encoded in a proprietary chart, the portable asset is the mapping, not the vendor’s schema. The semantic layer outlives any single ERP.
- **ERP migration as a first-class use case.** Migrating from ERP A to ERP B is, in Kontablo’s terms, an import (ERP A tree $\rightarrow$ graph) followed by an export (graph $\rightarrow$ ERP B tree). The universal pivot is the migration bridge: rather than building a bespoke $A \rightarrow B$ chart crosswalk by hand, both charts map once to the shared node set, and historical provenance is preserved through the UUID lineage.

In short, homogenization answers the Babel problem by abolishing heterogeneity; Kontablo answers it by making heterogeneity interoperable, which is the only option available to the large majority of enterprises that will never run on a single ERP instance.

3 The Downstream Economic Impact of Ontology Bridges

3.1 Financial Integration Velocity (FIV) in M&A

In corporate Mergers and Acquisitions (M&A), the speed at which the acquirer can unify the general ledgers of disparate subsidiaries is a primary constraint on synergy capture. Standard ERP-level integration requires manual mapping of charts of accounts, a process commonly reported at 90 to 180 days of consultancy labor. During this period, the parent company operates with restricted financial visibility, delaying cost-rationalization and capital-allocation decisions.

We propose **Financial Integration Velocity** (FIV) as a heuristic metric for this process — a stipulative definition we introduce, not an empirical result or a theorem — defined as the rate of successful account unification:

$$FIV = \frac{\sum_{i=1}^n a_i(1 - \epsilon_i)}{t_{\text{reconciliation}}} \quad (1)$$

where a_i represents the active account nodes in the target ledger, ϵ_i the error rate of the automated mappings, and $t_{\text{reconciliation}}$ the time in hours required to achieve validation. Under traditional manual methodologies, $t_{\text{reconciliation}}$ scales with the complexity of the chart of accounts. By deploying the Level 3 mapping engine as a translation layer, Kontablo is *designed* to drive $t_{\text{reconciliation}}$ toward a small constant — a design target on the order of hours rather than months, since only Tier 3 low-confidence mappings require human-in-the-loop review. We present the sub-24-hour figure as a target the architecture is built to enable, not a measured outcome; quantifying realized $t_{\text{reconciliation}}$ on live M&A integrations is explicit future work.

3.2 An Illustrative Cost Scenario for SME Operating Margins

Small and Medium Enterprises (SMEs) face disproportionate accounting costs when expanding across borders. Legacy ERP vendors price multi-entity consolidation modules out of reach for growing firms, forcing them to rely on outsourced accountants. By deploying the Kontablo Level 3 bridge via lightweight API integrations (such as the microSaaS prototype), SMEs can automatically map local tax invoices to a global standard.

As an *illustrative scenario* — a hypothesis to be validated with real deployment data, not an empirical result — consider an export-oriented SME with annual revenues of \$5M whose cross-border translation overhead and outsourced bookkeeping costs are substantially automated by such a bridge. If automation reduced those specific costs by 70–85% (the plausible upper range for the routine, deterministic share of the work, consistent with the volume-coverage structure described in Section 13.4), the saving would translate to a net operating margin expansion on the order of 1.8–3.2 percentage points. We present this as a motivating model with stated assumptions; measuring realized savings in actual SME deployments is explicit future work.

3.3 Real-Time Auditing (RTA) vs. Retrospective Audits

The traditional corporate audit is a retrospective, batch-processed event. Auditing firms spend months sampling journal entries from the preceding year, looking for anomalies after they have already affected the balance sheet. This latency permits fraud to remain undetected for long periods and distorts financial reporting.

We compare the latency of anomaly detection. Under a standard retrospective audit, the audit latency (L_{audit}) is the gap between when a transaction is posted and when the audit that examines it is reported:

$$L_{\text{audit}} = T_{\text{report}} - T_{\text{transaction}} \approx 30 \text{ to } 365 \text{ days} \quad (2)$$

With Kontablo’s Co-responsibility Architecture (CRA), the deterministic boundary checks execute at the transaction level. If an AI agent or accountant commits a mapping that violates a deterministic boundary, an inconsistency flag is raised and logged at write time, so the detection latency for that class of anomaly (L_{rta}) is:

$$L_{\text{rta}} = t_{\text{execution}} + t_{\text{flag}} \approx < 1 \text{ second} \quad (3)$$

RTA complements, it does not replace, the existing audit. We are careful not to overclaim here: real-time boundary enforcement does not make the forensic auditor obsolete. It adds three capabilities on top of the existing method rather than substituting for it. (1) *Real-time flagging*: boundary-violating mappings are caught at the instant of posting (Eq. 3) instead of months later (Eq. 2), but this reaches only the anomalies a *deterministic* rule can express — nature, liquidity, statement-class, and the other Boundary-Library invariants — not judgement-level fraud or collusion. (2) *Queryable, AI-tagged lineage*: because every posting carries a universal ID, a confidence score, and the AI’s justification, an auditor can query the population directly (“show every override against a deterministic rule,” “every Tier-3 mapping below 0.7”) rather than sampling blind. (3) *Traditional sampling remains valid and is augmented*: random-sample forensic testing, analytical review, and professional skepticism continue unchanged — now applied to a cleaner, fully-traced population with the routine boundary errors already removed. The net effect is an *increase* in forensic auditability and capacity, not a replacement of the auditor: the role shifts from detecting mechanical mis-postings by hand toward investigating the higher-order anomalies the deterministic layer cannot adjudicate.

3.4 Auditability and Fraud Prevention

Manual reconciliation is inherently opaque: when accounting teams manually adjust entries in spreadsheets during month-end close, they create an audit gap between the source ledger and the consolidated statement.

Kontablo’s contribution to auditability is structural rather than statistical: every transaction recorded via the Kontablo Graph is linked to a deterministic universal ID, so a forensic audit can trace any consolidated line item back to its origin in the local ledger (e.g., a Vietnamese VAS account) with no semantic translation lost in the process, and every human or AI override is permanently flagged in the audit trail (Section 9). We deliberately do not quantify this as a percentage improvement — auditability gains depend on the baseline process being replaced — but the mechanism replaces sample-based retrospective auditing with complete, queryable lineage. For an AI-driven economy, this translates to **Real-Time Auditing (RTA)** rather than once-a-year forensic sampling.

3.5 Foreign-Currency Reconciliation and the Transaction-Date Rate (IAS 21)

The hyperinflation cases of Section 13 exercise IAS 29 dual-rate divergence, but a more pervasive foreign-currency problem sits one level below it and applies to *every* cross-border ledger, not only hyperinflationary ones: IAS 21 requires a foreign-currency transaction to be recorded at the spot exchange rate *at the date of the transaction*. In practice, when transactions are not reconciled as they occur, firms fall back on the standard’s practical expedient — a period average or weighted-average rate applied at month-end close. That expedient is the quiet source of two distinct problems. First, it is an *approximation*: an average rate applied to a month of transactions silently misstates any individual transaction whose true spot rate diverged from the mean, and the error grows with both volatility and the dispersion of transaction timing within the period. Second, and worse,

batch period-rate translation is a *cover*: because the reconciliation happens away from the transaction and away from its date-stamped rate, an improperly classified or timed entry is far harder to detect once it has been folded into an averaged month-end figure.

A common semantic layer addresses both at their root by making the rate a property of the transaction rather than of the period. Each posting carries its transaction-date spot rate alongside its universal ID and provenance, so month-end translation becomes a deterministic re-aggregation of date-stamped amounts rather than the application of a single smoothing rate, and any entry can be re-derived to its exact transaction-date rate on audit. This does not repeal the practical expedient — a firm may still elect average-rate reporting where IAS 21 permits it — but it converts the expedient from a lossy necessity into an explicit, reversible choice, and it removes the reconciliation blind spot in which mis-timed or mis-rated entries currently hide. Implementing per-transaction rate capture and the IAS 21 average-rate election as a first-class ledger dimension is on the roadmap; the architecture already carries the per-transaction provenance the mechanism requires.

4 The Kontablo Ontology: Bridging the Divide

4.1 Graph Architecture vs. Rigid Trees

Kontablo’s fundamental departure from existing commercial ERP parameters is its graph-based metadata model. Extant commercial and statutory systems (NetSuite, Mexico’s SAT, France’s PCG) use rigid tree hierarchies which fundamentally fail at massive multi-jurisdiction mappings.

When a company attempts to map its Mexican subsidiary (with 5-digit SAT codes) to its French holding company (with 3-digit PCG codes) within a traditional ERP, it encounters a 1:N or N:M mapping problem that can only be resolved via labor-intensive manual overrides. Figure 2 illustrates this structural limitation: in any statutory tree, each account node has exactly one parent, so an account cannot simultaneously satisfy a tax-ledger hierarchy, a cost-center hierarchy, and an IFRS balance-sheet hierarchy.

Kontablo resolves this by abstracting the transaction outside of a local country’s tree; it becomes an incoming edge to a **Universal Graph Node** which serves as the absolute pivot point. Figure 3 shows how local nodes from multiple jurisdictions converge on a single universal node via deterministic lookup or AI semantic fallback, and how that node maps forward to the IFRS taxonomy.

5 The Ontological Bridge: Tree to Graph

5.1 The Graph-Based Metadata Model

Traditional accounting ledger structures are restricted by the relational tree paradigm: each account belongs to exactly one parent group (e.g., *Petty Cash* → *Cash* and *Cash Equivalents* → *Current Assets* → *Assets*). This rigid hierarchy prevents accounts from expressing multi-dimensional membership, such as being classified simultaneously under a statutory tax ledger, a functional cost center, and a consolidated IFRS report.

Kontablo solves this by representing the accounting universe as a directed multi-dimensional property graph. Formally, we define the accounting ledger space as a graph G :




Figure 2: Example of a rigid statutory chart-of-accounts tree (Mexico SAT excerpt). Each account node has exactly one parent. This single-path constraint prevents multi-dimensional classification: an account cannot simultaneously satisfy a tax-ledger hierarchy, a cost-center hierarchy, and an IFRS balance-sheet hierarchy — the mapping problem that Kontablo’s graph architecture resolves.

$$G = (V, E, \lambda, \mu)$$

Where:

- $V = V_{\text{local}} \cup V_{\text{universal}}$ is the set of nodes, representing localized statutory chart-of-accounts (COA) nodes (V_{local}) and standardized Level 3 Kontablo universal nodes ($V_{\text{universal}}$).
- $E \subseteq (V \times V)$ is the set of directed edges. Edges represent either transactional flows (debit/credit postings between accounts) or semantic mapping relationships (linking a localized account node to a universal target node).
- λ is a node-labeling function assigning each vertex a tuple of analytical dimensions in the product of their domains,

$$\lambda : V \rightarrow \mathcal{U} \times \mathcal{N} \times \mathcal{S} \times \mathcal{F}_{\text{fis}} \times \mathcal{F}_{\text{fun}} \times \mathcal{K} \times 2^{\mathcal{J}},$$

where each factor is declared explicitly — the point being that *every node dimension is non-numerical*: a finite categorical set, an identifier set, or a power set, compared by equality or membership, never by arithmetic.

- $\mathcal{U}$ — the UUID identifier set (128-bit RFC 4122 labels); a categorical name space, not $\mathbb{N}$.
- $\mathcal{N} = \{\text{debit}, \text{credit}\}$ — nature; a 2-element set.
- $\mathcal{S} = \{\text{balance_sheet}, \text{income_statement}, \text{cash_flow}\}$ — statement class; a 3-element set.
- $\mathcal{F}_{\text{fis}}$ — a finite set of fiscal/tax tags (VAT direction, withholding, e-invoice tagging, Zakat).
- $\mathcal{F}_{\text{fun}}$ — a finite set of functional / cost-center labels.




Figure 3: The Kontablo graph architecture: local statutory account nodes (Mexico SAT, France PCG, Vietnam VAS) converge on a universal Kontablo node via deterministic exact lookup or AI semantic fallback. The universal node maps 1:1 to the IFRS taxonomy target. Autonomous AI agents interact directly with the universal layer via the AP2 protocol.

- $\mathcal{K} = \mathcal{E} \times \{\text{eliminate, retain}\}$ — the consolidation role: counterparty entity $\mathcal{E}$ and elimination flag.
- $2^{\mathcal{J}}$ — the power set of the 195 jurisdictions $\mathcal{J}$; overlay membership is a *subset*, not a scalar.

These categorical domains are precisely why the deterministic boundary checks (Section 8) are equality and membership tests rather than numeric comparisons.

- $\mu : E \rightarrow \mathcal{P}$ is an edge-property mapping function, and it is on *edges*, not nodes, that the only numeric quantities live. For transactional edges, $\mathcal{P} = \mathcal{G} \times \mathcal{T} \times \mathcal{U}$: an amount in the Pacioli group $\mathcal{G} \cong \mathbb{R}$ (signed monetary value, currency-tagged and FX-normalized to a numéraire; a fixed-point rational $\subset \mathbb{Q}$ in implementation, Section 6.5); a timestamp in a totally-ordered time domain $\mathcal{T} \cong \mathbb{R}_{\geq 0}$ (an integer epoch $\subset \mathbb{N}$ in implementation); and an initiating_Agent_ID $\in \mathcal{U}$ (categorical). For mapping edges, $\mathcal{P} = [0, 1] \times \mathcal{M} \times \mathbb{B}$: a confidence_score $\theta \in [0, 1] \subset \mathbb{R}$ (a bounded, totally-ordered, thresholded heuristic score — *not* a measure-theoretic probability); a match_method $\in \mathcal{M} = \{\text{tier1_exact, tier2_keyword, tier3_semantic}\}$; and audit_flags $\in \mathbb{B} = \{\top, \perp\}$ (boolean). No dimension ranges over $\mathbb{C}$, and signed magnitude is carried by the Pacioli group rather than by an $\mathbb{R}_{\geq 0}$ magnitude plus a detached sign.

What the dimensions mean. Set-theoretic notation aside, the point for a practitioner is concrete: each universal node carries, *simultaneously*, the full set of analytical attributes that a conventional chart of accounts would otherwise scatter across separate hierarchies. Figure 4 renders one such node — a trade receivable arising from a credit sale (`asset.current.receivables`) — with every dimension attached at once:

- **Statement classification** — where the line lands in the financial statements: balance sheet, profit & loss, or cash flow. Our receivable is a balance-sheet (current

asset) item.

- **Nature** — whether the account’s normal balance is a debit or a credit. The receivable is debit-natured.
- **Fiscal / tax** — the tax treatment the line carries: VAT direction (input vs. output), withholding obligations, electronic-invoice tagging (Mexico’s SAT, Brazil’s SPED), or Zakat assessment in Gulf jurisdictions.
- **Functional / cost-center** — the managerial dimension: the department, project, or cost center to which the amount is attributed for internal reporting.
- **Consolidation / intercompany** — the group-reporting role: the counterparty entity and whether the balance must be eliminated on consolidation (an intercompany receivable nets against the matching payable in the other entity’s books).
- **Jurisdiction overlay** — which local statutory code(s) map into this universal node; e.g. MX-SAT-105 (Mexico), FR-PCG-411 (France), and VN-VAS-131 (Vietnam) all resolve to the same universal receivable.

In a single-parent *tree*, each of these would demand its own hierarchy — one for the tax ledger, one for the cost center, one for the IFRS balance sheet, one for consolidation — and the same account would have to be *duplicated* under each, with manual reconciliation keeping the copies in step. The graph keeps one canonical node and lets the dimensions coexist as properties and edges on that node, so a change made to the receivable is made once and stays consistent across every view.

5.2 Level 3 Minimum Core Taxonomy

The core of the Kontablo protocol is the **Level 3 account specification**, which maps the minimum viable set of accounts required to express the financial position of any enterprise, regardless of its industry or country. Level 3 accounts act as universal semantic anchors in $V_{\text{universal}}$.

The taxonomy defines exactly 30 Level 3 accounts, divided into two distinct components:

1. **18 Near-Universal Core Accounts:** concepts that recur near-universally across the committed jurisdiction mapping sets (the most universal appear in 94–98% of them; none reaches a literal 100%). These cover core operational categories including Cash, Trade Receivables, Trade Payables, Retained Earnings, Cost of Goods Sold, and Administrative Expenses.
2. **12 Optional/Jurisdictional Overlays:** Activatable nodes based on specific regulatory regimes or business models. Examples include:
 - `asset.current.vat_input` and `liability.current.vat_output` (disabled in jurisdictions without VAT, such as the United States).
 - `liability.current.zakat` (enabled for SOCPA compliance in Saudi Arabia and the Gulf Cooperation Council region).
 - `asset.noncurrent.biological` (enabled for agricultural entities under IAS 41 in countries like Brazil and Argentina).




Figure 4: A single universal Kontablo node carries every analytical dimension at once. One node — a trade receivable from a credit sale (**asset.current.receivables**) — simultaneously holds its statement classification, debit/credit nature, fiscal/tax treatment, functional cost-center attribution, consolidation/intercompany role, and the local statutory codes (Mexico SAT, France PCG, Vietnam VAS) that map into it. In a single-parent tree each of these dimensions would require a separate hierarchy and a duplicated copy of the account; the graph keeps one canonical node and attaches the dimensions as coexisting properties and edges.




Figure 5: The Kontablo mapping bridge connects localized statutory account nodes to standardized IFRS classifications via a weighted edge that carries a confidence score and match method.

A reproducible transaction-frequency benchmark (Section 13.4; `scripts/coverage_benchmark.py`) estimates that this 30-account configuration captures **~94% of routine transaction volume** by posting count in a representative SME ledger profile, with a 34-account extended core reaching **~99%**.

5.3 Tree-to-Graph Compatibility Protocol

All major legacy ERPs (ERPNext, Odoo, Zoho Books, SAP B1) enforce strict hierarchical trees where each account has exactly one parent. Integrating these systems with Kontablo requires managing the mismatch between trees and graphs. Kontablo implements a two-way compatibility protocol (ADR-008²):

5.3.1 Tree-to-Graph (Import Stream)

When importing ledger data from a legacy ERP into Kontablo, the linear account hierarchy is converted into a multi-dimensional graph. The mapping service traverses the ERP tree and assigns each leaf node a mapping edge pointing to a Level 3 universal node in Kontablo, enriching the node with its multi-dimensional properties:

$$\text{ERP Leaf} \xrightarrow{\text{Mapping Edge}} \text{Kontablo Node } V_u \quad (4)$$

If the import tool encounters a group account (e.g., an asset class summary node), it collapses the group and maps only the underlying transaction-bearing leaf nodes, preserving historical ledger provenance.

Worked example. An ERPNext chart contains the leaf **1310 - Debtors** under the parent path **Application of Funds → Current Assets → Accounts Receivable**. The import stream discards the single-parent path and attaches one mapping edge **1310 - Debtors** $\xrightarrow{\text{maps-to}}$ **asset.current.receivables**, then enriches the node with the dimensions the tree could not express simultaneously — IFRS tag (**IAS 1.54(h)**), nature (**debit**), statement (**balance_sheet**), and any jurisdiction overlay — none of which required a separate parent in the source tree.

5.3.2 Graph-to-Tree (Export Stream / Linearization)

To export data back to a legacy ERP, Kontablo must **linearize** the multi-dimensional graph back into a single-parent tree. The linearization algorithm selects the primary reporting dimension (typically the IFRS Balance Sheet classification edge) and ignores secondary analytical dimensions. This maintains downward compatibility with ERP

²Architecture Decision Records (ADRs) referenced throughout this paper are committed in the project repository under `docs/adr/`; ADR-008 specifies the tree-to-graph compatibility protocol, ADR-009 the determinism-over-stochasticity principle, and ADR-011 per-client determinism.

database schemas while keeping the multi-dimensional graph intact inside Kontablo for analysis.

Worked example. Exporting the same `asset.current.receivables` node into an Odoo chart, the linearizer picks the balance-sheet edge as the primary dimension and emits a single-parent account `Current Assets → Receivables → Trade Receivables`; the node’s secondary analytical edges (e.g. a cash-flow-statement classification or an intercompany-counterparty dimension) are dropped from the exported tree but retained inside Kontablo. The discarded secondary edges are exactly the information a single-parent tree cannot hold — the formal source of the reconciliation obstruction analysed in Section 6.4.

5.4 AI Validation and the Co-responsibility Architecture

Kontablo maps 1:N cardinality disruptions via a strictly defined **Co-responsibility Architecture**. In this model, the integrated large language model (LLM) acts as an augmentative auditor that proposes mappings and verifies human overrides; the architecture is model-agnostic by design and names no specific vendor.

If a human accountant makes a selection that contradicts the deterministic boundary of a node (e.g., mapping a physical cash account to a non-current asset node), the system generates an **Inconsistency Flag**. This flag, accompanied by a machine-generated audit note, is permanently attached to the transaction. This ensures that while the human remains the final authority (“AI proposes, Human disposes”), the system enforces a digital trail of accountability that simplifies future forensic audits. Final legal accountability rests with the human operator. The AI’s role is verification and warning; the human retains veto power and bears sole legal responsibility for any mapping committed to the ledger. The audit trail records both the AI’s proposal and the human’s disposition, making every decision traceable and attributable.

5.5 Standard Hierarchy: The Level 3 Minimum Core

Based on a comparative analysis of 195 jurisdictions, we identify 18 Level 3 accounts that recur near-universally across the analyzed standards (e.g., cash, receivables, trade payables, share capital): in our committed jurisdiction mapping sets, the most universal of these concepts appear in 94–98% of jurisdictions, with the remainder reflecting incomplete mapping depth rather than conceptual absence from the underlying standard. An additional 12 accounts (e.g., Input VAT, Zakat Provision, Deferred Tax) are classified as “Jurisdiction-Specific Overlays,” which are toggled based on the entity’s regulatory environment metadata.

6 Mathematical Foundations of the Translation Mechanism

The tree-to-graph-to-tree pipeline of Section 5 is not an ad hoc engineering convenience; it is an instance of a construction that recurs, under different names, across category theory, applied topology, and computational linguistics. Making the construction explicit serves two purposes: it shows *why* the round trip resolves the N:M mappings that defeat direct tree-to-tree approaches, and it makes the frontier of the method legible — it tells

us exactly where deterministic guarantees hold and where irreducible reconciliation loss begins.

6.1 The Babel Problem, Stated Formally

Let a localized statutory chart of accounts be a rooted tree T_s and a target statutory chart be a rooted tree T_t , each carrying the single-parent constraint of Section 5. The “Accounting Babel” is the observation that there is in general *no faithful structure-preserving map* $T_s \rightarrow T_t$: an account in T_s may correspond to several accounts in T_t and vice versa (the $1:N$ and $N:M$ cardinalities), and the multi-dimensional roles an account plays (statutory, fiscal, functional, consolidated) cannot be expressed inside a single-parent hierarchy at all. The direct approach fails because the two trees do not share a common coordinate system.

A concrete instance. Consider one economic event — a domestic credit sale of goods with value-added tax. In Mexico’s SAT chart it touches 105 (Clientes), 401 (Ingresos), and 208 (IVA trasladado); in France’s PCG, 411 (Clients), 707 (Ventes de marchandises), and 44571 (TVA collectée); in Vietnam’s VAS, 131 (trade receivables), 511 (revenue), and 3331 (VAT payable). No single-parent edge relates “SAT 208” to “PCG 44571” to “VAS 3331”: they are the *same* concept (output VAT payable) under three incompatible coordinate systems, and the receivable line additionally carries fiscal and functional roles the tree cannot hold at once. This is the $N:M$ obstruction the formalism names — and the reason the pivot below routes every chart through one shared node set rather than attempting $T_s \rightarrow T_t$ directly.

6.2 Lift, Pivot, Project: an Adjoint Pair

Kontablo resolves this exactly as machine translation resolves the analogous problem between natural languages: rather than mapping T_s directly to T_t , it analyses T_s up into a jurisdiction-independent intermediate representation and then generates T_t down from it. In machine translation this intermediate is the *interlingua* at the apex of the Vauquois triangle (Vauquois, 1968); in Kontablo it is the Level 3 universal node set $V_{\text{universal}}$, which serves as the absolute pivot.

This pivot construction has a natural categorical reading. Following Spivak (2012), a database schema is a small category and an instance is a set-valued functor on it; a schema mapping $F : \mathcal{S} \rightarrow \mathcal{T}$ then induces three canonical *data-migration functors*. Kontablo’s import and export streams are structurally analogous to two of them:

$$\Sigma_F : \mathcal{S}\text{-Inst} \rightarrow \mathcal{T}\text{-Inst} \quad (\text{import / lift: tree} \rightarrow \text{graph})$$

$$\Delta_F : \mathcal{T}\text{-Inst} \rightarrow \mathcal{S}\text{-Inst} \quad (\text{export / project: graph} \rightarrow \text{tree})$$

where Σ_F is left adjoint to Δ_F . The import stream enriches each leaf of the local tree with a mapping edge to a universal node and lifts it into the multi-dimensional graph G ; the export stream linearizes G back to a single-parent tree. We state the correspondence as an analogy, in the same register as the Cover’s-theorem passage below, because the full construction has not been carried out: we have not defined the schema categories, constructed F , or verified the adjunction, and the implemented export step *selects* a primary reporting dimension (Section 5) — a choice a canonical adjoint would have to eliminate or parameterize. Working the correspondence out for a minimal chart pair, or amending it where it fails, is labeled future work; what the analogy buys today is a design

discipline (imports should be free constructions, exports should be structure-forgetting) rather than a theorem.

6.3 Why Lifting the Dimension Resolves the Mapping

The intuition that a problem unsolvable in a low-dimensional structure becomes solvable once lifted into a richer one is formalized, in a different setting, by Cover’s theorem (Cover, 1965): a classification that is not separable in a low-dimensional space becomes separable, with high probability, once mapped nonlinearly into a higher-dimensional space. We invoke this only as an analogy of *principle* — Cover’s setting is geometric separability, not structure-preserving translation — but the shared moral is exact: the $N:M$ entanglement that cannot be resolved inside a single-parent tree is resolved once the account is re-expressed as an edge into a multi-dimensional graph, where membership in several analytical dimensions is no longer in conflict.

6.4 The Cost of the Return Trip: a Reconciliation Obstruction

The return trip is lossy by construction. Linearization selects a primary reporting dimension and discards the secondary analytical edges (Section 5); the discarded content is precisely the failure of the graph G to be a tree. The natural question — *when can the local ledgers of many jurisdictions be glued into one globally consistent account, and what obstructs it?* — is the question that sheaf cohomology was built to answer.

Assign to the mapping graph a cellular sheaf whose stalks carry the local ledger data and whose restriction maps encode the consistency conditions on shared (overlapping) accounts (Robinson, 2017; Curry, 2014; Hansen & Ghrist, 2019). The domains here are explicit: each stalk is a finite-dimensional real vector space $\mathbb{R}^d$ (the monetary amounts a node carries across its d analytical dimensions, valued in the Pacoli group of Section 6.5), the restriction maps are $\mathbb{R}$ -linear, and the cohomology spaces are therefore $\mathbb{R}$ -vector spaces. Then the degree-zero cohomology H^0 is the space of *globally consistent reconciliations* — the data that patches together across all jurisdictions — while a non-trivial degree-one class, $H^1 \neq 0$, is the exact obstruction to such gluing: it measures, and localizes, the inconsistency that no choice of mapping can remove.

This gives two of Kontablo’s operational notions an intended formal referent. A node’s *deterministic boundary* (ADR 009) is designed to play the role of a sheaf restriction map: it constrains which local sections are admissible. The *Inconsistency Flag* of the co-responsibility architecture is designed to play the role of a witness to a local section that fails to extend — a non-zero cohomology class. The implemented flag is, today, a set of deterministic rule checks (Section 13.5), not an instantiated sheaf; the current state of formal validation is a committed toy experiment (`research/experiments/kirchhoff_hodge_toy/`) that constructs a cellular sheaf over a small intercompany scenario, computes δ^0 and the H^0/H^1 dimensions numerically, and resolves the imbalance by Moore–Penrose reconciliation — a mechanics validation, not a proof of the general correspondence. Determinism and reconciliation are nonetheless designed to speak one language: the deterministic scaffolding defines the restriction maps, and cohomology is the intended formal account of where global consistency fails and a human must dispose.

6.5 The Accounting-Native Anchor: the Pacioli Group

None of this is foreign to accounting’s own mathematics. Ellerman (2014) shows that double-entry bookkeeping is, exactly, the *group of differences* construction (the “Pacioli group”) on unsigned T-account pairs, and — crucially for Kontablo — that this construction *generalizes naturally to multi-dimensional vector accounting*, in which scalar money amounts are replaced by vectors recording distinct types of property or analytical dimension. This fixes the domain of monetary value precisely. The scalar Pacioli group is the group of differences of ordered pairs of non-negative amounts, $\mathcal{G} = (\mathbb{R}_{\geq 0} \times \mathbb{R}_{\geq 0}) / \sim \cong (\mathbb{R}, +)$: a signed real, where the debit/credit nature supplies the sign rather than a separate $\mathbb{R}_{\geq 0}$ magnitude. The multi-dimensional generalization is $\mathcal{G}^d \cong \mathbb{R}^d$ over the d analytical dimensions — the domain the sheaf stalks of Section 6.4 are valued in. (Implementations realize $\mathcal{G}$ as fixed-point rationals $\subset \mathbb{Q}$ for exact cent arithmetic; the algebra is stated over $\mathbb{R}$.) Kontablo’s multi-dimensional property graph is the data-structure realization of precisely that generalization. The contribution is therefore not to invent a new algebra of accounting, but to operationalize, as a deterministic protocol over a property graph, a multi-dimensional generalization that has been latent in the mathematics of double-entry bookkeeping since its first formal treatment.

7 Related Problems and Convergent Prior Art

The translation mechanism of the previous section is not unique to accounting. Several mature fields solve the same problem — reconcile two structures that share no common coordinate system — by the same move: lift to a richer intermediate structure, operate there, project back. We position Kontablo against this prior art at three explicit strengths of claim, so that a rigorous instance is never mistaken for a suggestive metaphor.

7.1 Foundation: constructions Kontablo instantiates

Three bodies of work supply the actual machinery and are cited as ground truth: functorial data migration (Spivak, 2012) and ologs (Spivak & Kent, 2012) for the schema-to-schema translation; cellular-sheaf cohomology for data fusion (Robinson, 2017; Curry, 2014; Hansen & Ghrist, 2019) for the local-to-global reconciliation obstruction; and the Pacioli group with its multi-dimensional generalization (Ellerman, 2014) for the accounting-native algebra. These are not analogies — Kontablo’s import/export and reconciliation are instances of them.

7.2 Sibling: fields solving the same problem the same way

Ontology and schema matching in the semantic web (Thiéblin et al., 2020) target precisely the heterogeneity Kontablo calls Babel, and the Ontology Alignment Evaluation Initiative provides standardized benchmarks worth borrowing for evaluation. The most literal twin is *phylogenetic reconciliation*: a gene tree and a species tree, distinct in topology, are reconciled by embedding one into the other and explaining the discordance by discrete events (duplication, loss, transfer), with a rich complexity theory of when the reconciliation is tractable (e.g., Hasić & Tannier, 2019). This is the tree-to-tree-via-richer-structure problem with a different vocabulary, and its cost models are a direct precedent for

quantifying linearization loss. Interlingua machine translation (Vauquois, 1968) supplies the pivot architecture itself.

7.3 Machine learning: structural alignment as optimization

A direct line of machine-learning work attacks the same alignment problem and is the most likely source of practical leverage. Gromov–Wasserstein optimal transport (Mémoli, 2011; Xu et al., 2019) aligns two relational structures that live in *incomparable* spaces by matching their internal distance structure alone, without any shared feature coordinates — exactly the situation of two charts of accounts with no common key. It yields a coupling between the nodes of a source structure and a target structure as the solution of a well-posed optimization problem.

Its role in Kontablo is deliberately bounded and subordinate to ADR 009. Gromov–Wasserstein alignment is a candidate for *mapping-proposal generation under sparse labels*: an optimization-based, reproducible alternative to stochastic language-model guessing when no deterministic rule yet exists for a given account. It is never the mapping authority. The optimal coupling is a proposal that must pass the deterministic boundary and human disposition; production mapping remains deterministic and auditable.

7.4 Neurosymbolic AI and Ontology-Constrained Enterprise Agents

A growing body of literature on neurosymbolic AI — integration of formal symbolic reasoning with neural language models — directly converges on the design principles underlying Kontablo’s harness architecture. We position Kontablo within this literature at the level of application and contribution.

Ontological grounding reduces financial hallucination empirically. Luong Tuan & Sanyal (2026) present a controlled experiment (1,800 runs, five regulated industries, three LLMs) demonstrating that ontology-coupled agents significantly outperform ungrounded baselines on metric accuracy and role consistency. Independently, [arXiv:2603.04663] (2026) introduces a *Universal Fact Ledger* — deterministic typed variables grounding financial agent reasoning — and shows this compresses financial hallucination to 1.2%, with residual error concentrating at domain boundaries. Both findings are direct empirical support for Kontablo’s architectural Principle 5 (determinism over stochasticity wherever movable).

Ontologies provide categorically different guarantees than RAG. While retrieval-augmented generation (RAG) improves probabilistic accuracy, it does not provide structural constraints on the output space. Sharma et al. (OG-RAG, EMNLP 2025) demonstrate +55% fact recall from ontology-anchored hypergraph retrieval, but acknowledge that formal constraint enforcement requires machine-readable ontological definitions that document-chunk retrieval cannot supply. This distinction is central to Kontablo: the UUID graph membership check is not a probabilistic improvement over RAG — it is a categorical rejection mechanism applied before any value reaches the ledger.

Kontablo’s contribution relative to this literature. Existing neurosymbolic enterprise architectures (Luong Tuan & Sanyal, 2026; Foundation AgenticOS) address general enterprise domains (fintech, healthcare, insurance) with industry-specific ontologies. Kontablo specializes this architecture to *universal accounting* with three contributions

not present in the general case: (1) a UUID-keyed graph that is jurisdictionally composable across all 195 sovereign states, (2) deterministic boundary predicates that enforce accounting-native invariants (debit/credit nature, IAS 29 restatement rules, Zakat base exclusions) independently of the language model, and (3) an explicit universality proof via the IFRS-pure exhaustion argument (Appendix D) that the coverage boundary is finite and structured, not open-ended.

7.5 Open accounting problems the same mechanism addresses

Framing the mechanism formally widens its surface of application. Three open problems in global accounting reduce to the same lift–reconcile–project pattern. **Consolidation and intercompany elimination** is a gluing of local ledgers into a global one — the precise scenario for which the sheaf-cohomological obstruction H^1 was introduced, so a failed consolidation has a computable, localizable cause. **Multi-GAAP reconciliation** (IFRS $\leftrightarrow$ US GAAP $\leftrightarrow$ local) is span completion through the universal pivot, with the residual obstruction marking which reconciliations are impossible without loss. **Chart-of-accounts versioning and provenance** is functorial migration across time: a schema mapping between an old and a new chart induces the migration functors that move historical data and queries consistently (Spivak, 2012). Each is a candidate spin-out result rather than a claim the present paper proves.

8 Deterministic Logic Rules vs. AI Probability

8.1 The Failure of Pure Semantic AI in Accounting

While modern Large Language Models (LLMs) excel at semantic matching, using them as primary systems of record in financial systems is dangerous. LLM hallucinations in bookkeeping can lead to tax fraud, regulatory penalties, and audit failures. A mapping accuracy of 95% is acceptable for search queries, but unacceptable for general ledgers where accuracy must be 100% (or at least all discrepancies must be flagged). An AI proposing that an electricity bill is a “Prepaid Expense” instead of an “Administrative Expense” introduces systemic errors that compromise financial integrity.

8.2 The Three-Tier Resolution Strategy

To achieve high automation without risking hallucinated mappings, Kontablo implements a **Three-Tier Resolution Strategy** that combines deterministic rules with probabilistic AI fallbacks. When an unmapped localized ERP account is processed, the system routes it through the following pipeline:

1. **Tier 1: Exact Lookup (Direct Mapping):** The engine queries the central Level 3 YAML schema. If the local account code matches a pre-mapped standard code for that jurisdiction (e.g., Mexican SAT code 101 matches `asset.current.cash` directly), a link is created. This tier produces an exact match with a confidence score of 1.0 and requires zero inference latency.
2. **Tier 2: Disambiguation Rules (Rule-Based Regex):** If Tier 1 fails, the engine applies keyword and regular expression rules to the account type and name. For

example, a Zoho Books account of type `other_current_liability` with a name containing “VAT” or “IVA” resolves to `liability.current.vat_output` (confidence 0.80–0.90).

3. **Tier 3: Semantic AI Fallback (Probabilistic Matching):** If both Tier 1 and Tier 2 fail, the engine activates the Semantic Matcher AI. The engine generates a vector embedding of the account’s name, type, and description, comparing it against the embeddings of the 30 Level 3 universal accounts. The model outputs a mapping proposal along with a confidence score (0.50 to 0.85).

Figure 7 traces the same logic on three concrete inputs — one resolved at each tier — to make the distinction between the deterministic tiers and the single stochastic one concrete.

8.3 The Deterministic Questions Library

To prevent LLM mapping proposals from silently corrupting the ledger, all outputs from Tier 3 must pass through a strict library of deterministic boundary rules. Rather than letting the LLM decide a mapping in a single step, the system forces the mapping through a sequence of deterministic questions based on the target node’s metadata.

For example, if the AI proposes mapping a local account “Petty Cash Fund” to the Level 3 node `asset.current.cash`, the rule engine verifies:

1. **Nature Check:** Is the debit/credit nature of the local account consistent with the target node? (e.g. cash is debit).
2. **Liquidity Check:** Does the local account represent an asset available at par under 90 days?
3. **Boundary Check:** Is the local account a cash equivalent, restricted cash, or digital currency?

If any validation returns a mismatch (e.g., if the AI attempts to map a restricted-use cash account to a standard cash node, or maps a liquid account to a non-current asset node like PPE), the engine overrides the confidence score, setting it to a penalized value of 0.3, and raises an anomaly flag, forcing manual human-in-the-loop review.

8.4 Cross-Jurisdiction Semantic Bridge

Imagine a multinational with two accountants: one in Mexico who names the cash account “Caja Revolvente” and one in Brazil who uses “Caixa Geral.” To a human ERP manager, these are different. In the Kontablo Graph, both accountants’ local entries hit the same node because they both satisfy the deterministic criteria for ‘`asset.current.cash`’.

This allows a corporate headquarters to have a **Unified Single-target Ledger** where the name of the local account is irrelevant to the consolidated analytical truth. The consolidation occurs at the graph layer, not the string/label layer.




Figure 6: The Three-Tier Resolution Pipeline. Local ERP accounts pass through exact code checks (Tier 1) and rule-based regex parsing (Tier 2) before falling back to Semantic AI (Tier 3). All outputs enter the Co-responsibility Architecture (CRA): high-confidence proposals (≥ 0.80) are queued for pre-authorized review with audit logging; low-confidence or boundary-violating proposals require explicit human co-signature before ledger write-in. In no case does the agent commit a mapping to the ledger without a human-authorized or human-pre-authorized pathway.




Figure 7: The three tiers on concrete inputs. Tier 1 resolves a statutory code deterministically; Tier 2 resolves a name via an auditable multilingual rule when no code matches; Tier 3 (the only stochastic step) proposes a node when neither deterministic tier fires, and that proposal is constrained to existing universal nodes and gated by the Deterministic Boundary Library before it can post. The engine descends the tiers only on failure, so most volume resolves at Tier 1–2 without inference.

8.5 The Inverse Parametric Knowledge Effect: Why Non-Anglophone Jurisdictions Benefit Most

A non-obvious consequence of ontology-as-constraint is what Luong Tuan & Sanyal (2026) term the **Inverse Parametric Knowledge Effect (IPKE)**: the value of ontological grounding is *inversely proportional* to the LLM’s pretraining coverage of the target domain. For well-documented international concepts – Basel III ratios, PCI-DSS levels, IFRS core terminology – LLMs already encode sufficient parametric knowledge, and injecting additional ontological context produces marginal gains (or in some cases modest regressions due to context displacement). For enterprise-specific, non-English, or thin-corpus regulatory domains, ontological grounding is not an optimization but an operational requirement.

This effect has direct implications for Kontablo’s jurisdictional coverage strategy. The jurisdictions that benefit most from the deterministic constraint layer are precisely those where pretraining data is structutest sparse: VAS-Vietnam (government-mandated chart in Vietnamese), SOCPA-Saudi Arabia (Zakat provisions in Arabic), SYSCOHADA-francophone Africa (OHADA chart in French with local judicial overlays), and IAS 29 hyperinflationary contexts (Venezuela, Lebanon, Argentina) where the economic logic departs from standard IFRS in ways not well-represented in English-language corpora. The controlled experiment of Luong Tuan & Sanyal (2026) provides empirical support: Vietnamese-language financial domains showed 2× the improvement from ontological grounding relative to English-language domains ($\Delta = +.29$ vs. $+.12$ composite score), replicating across Claude Sonnet 4, Qwen 2.5 72B, and Gemma 4 26B.

8.6 Ontology vs. Retrieval-Augmented Generation

A natural alternative to a formal ontology is a Retrieval-Augmented Generation (RAG) pipeline over the same regulatory documents. Why is Kontablo not implemented as a

RAG system over IFRS standards, national tax codes, and SYSCOHADA blueprints?

RAG retrieves; it does not constrain. A RAG system can surface the relevant section of the Vietnamese VAS standard for a given transaction, but it cannot *prevent* the LLM from emitting an account code that does not exist in the client’s actual chart. The critical difference is the locus of the guarantee:

- **RAG** improves the probability that the model proposes a correct mapping. It does not eliminate invalid mappings – it makes them less likely.
- **Ontology-as-constraint** provides a categorical guarantee: any UUID that does not exist in the graph is rejected at the harness level before reaching the ledger, regardless of the model’s confidence.

Luong Tuan & Sanyal (2026) quantify the distinction: while RAG and ontology injection achieve comparable scores on lexical recall metrics, ontologies provide *categorically different* capabilities that document-chunk retrieval cannot replicate – hierarchical constraint enforcement, formal composability across industries, and machine-verifiable output validation (their Table 5). For Kontablo’s ledger-grade accuracy requirement – where a 95% accuracy rate is not acceptable and every mismatch must be flagged – the categorical guarantee is the operative property, not the average-case improvement. A system that *usually* maps correctly but *sometimes* silently posts to a non-existent account is not fit for financial purpose.

9 The Kontablo Agent: Harness Architecture and the Locus of Error

9.1 The Model–Harness Distinction

A common framing of AI reliability in critical systems asks: “*How do we prevent the LLM from hallucinating?*” This question, while intuitive, addresses the wrong architectural layer. It locates the reliability problem inside the model, where the design levers available to system architects are limited and where the state of the art offers no near-term guarantees suitable for ledger-grade accuracy.

A more productive framing asks: “*How do we design the system around the model so that the decision space the model can influence is constrained to valid outputs, and anything outside coverage escalates to a human rather than being confabulated?*” This reframe is the starting point of harness-centric agent design.

The **model** is the stochastic component: the Large Language Model that performs semantic interpretation, natural-language matching, and open-ended reasoning over unstructured input. The **harness** is the deterministic scaffold surrounding it: the tool definitions and schemas that structure its output, the ontological constraints that bound the valid output space, the validation layers that reject structurally invalid proposals before they reach the ledger, and the escalation protocols that route out-of-coverage cases to human review.

The core practical insight is that the reliability of an agentic system in a critical domain is determined far more by the harness than by the model’s intrinsic accuracy. This has consequences for where engineering effort should be concentrated and where

reliability guarantees can realistically be made. Recent work on agent scaffolding and tool-augmented language models substantiates this view. Willard & Louf (2023) demonstrate that grammar-constrained generation – where a formal grammar bounds the output token space – produces structurally valid outputs without sacrificing fluency. Luong Tuan & Sanyal (2026) extend this to enterprise multi-industry deployment across 1,800 controlled runs, showing that ontology-coupled agents significantly outperform ungrounded baselines on metric accuracy ($p < .001$, Kendall’s $W = .460$) and role consistency ($p < .001$, $W = .614$) across three model architectures. Critically, their neurosymbolic coupling maturity taxonomy formalizes the distinction between systems that constrain *inputs* (context injection, tool discovery, governance thresholds) and systems that validate *outputs* – a distinction we return to in Section 9.7.

9.2 Ontology-as-Constraint: Eliminating the Classical Accounting Hallucination Class

In the absence of a harness, an LLM deployed for transaction classification can hallucinate in an accounting-specific way: it may propose an account code or name that *sounds* plausible but does not correspond to any account in the client’s chart of accounts, or it may classify a transaction to the wrong node with high expressed confidence and no observable uncertainty signal.

The asymmetry with information retrieval is severe. A 95% classification accuracy is acceptable in search; in a general ledger, the 5% of misclassified transactions introduces systematic distortions that compound across reporting periods, survive into audited financial statements, and – in the presence of the Three-Tier Resolution pipeline – would do so *silently*, with no inconsistency flag raised, because an unconstrained model emitting a plausible-looking but wrong UUID would pass surface validation.

Kontablo addresses this through what we term **ontology-as-constraint**: the agent’s tool interface (exposed via the MCP server) accepts only Kontablo Level 3 UUIDs as valid account identifiers. The set of valid UUIDs is finite, enumerable, and maintained in the ontology graph. The harness rejects any proposed UUID that does not exist in the graph at validation time, before the value reaches the ledger write path.

This mechanism is analogous to grammar-constrained or schema-constrained generation in the structured output literature (Willard & Louf, 2023; Luo et al., 2025), where the valid output space is defined by a formal grammar and outputs violating that grammar are rejected at inference time. In Kontablo’s case the “grammar” is the ontology itself: the closed set of valid account UUIDs, each carrying its nature (debit/credit), IFRS tag, and deterministic boundary predicates from the Deterministic Boundary Library.

The consequence is that the *classical accounting hallucination class* – the emission of arbitrary account identifiers that do not exist in the ontology – is architecturally eliminated at the harness level. This is not a probabilistic reduction in hallucination rate; it is a categorical elimination of a specific error class within the constrained output space. The model can still be wrong within the set of valid UUIDs (mapping a transaction to `expense.admin` when `expense.cogs` is correct), but it cannot emit a non-existent account.

Implementation dependency: This guarantee holds only if the harness enforcement is strict and complete – no code path allows a UUID to reach the ledger write step without a graph-membership check. This is a design requirement, not an automatic consequence of using an LLM with tool definitions. Any deployment that bypasses the validation layer

for performance or convenience reasons loses this guarantee.

9.3 The Relocation of Residual Error

Eliminating the classical hallucination class does not eliminate error; it *relocates* it. The residual error population, after ontology-as-constraint is applied, concentrates at a specific and identifiable location: the **boundary of semantic coverage**.

Errors at the semantic coverage boundary take three forms:

1. **Unmapped transaction types.** A transaction type not yet represented in the ontology has no valid UUID. The Three-Tier Resolution Pipeline routes it to Tier 3 (semantic AI fallback); if confidence remains below the threshold (Section 6), it is flagged for human review rather than forcing a low-confidence posting. The error is not confabulated – it is escalated and made observable.
2. **Jurisdictional edge cases.** A jurisdictional rule not yet encoded in the Deterministic Boundary Library may allow a semantically plausible UUID to pass boundary validation while violating an implicit local constraint. These errors are attributable to incomplete rule coverage, not to model hallucination. They are also correctable through ordinary ontology maintenance as cases are discovered.
3. **Semantic ambiguity within coverage.** Some transactions satisfy the boundary predicates of more than one Level 3 node. In these cases Tier 3 may produce a defensible but suboptimal mapping. The Co-responsibility Architecture (Section 10) manages this residual by requiring human review for low-confidence Tier 3 outputs.

This relocation is architecturally significant because boundary errors are **observable**: they manifest as flagged inconsistencies, low-confidence outputs, and unmapped-transaction queues – all instrumentable signals. Stochastic interior errors from an unconstrained model are, by contrast, **silent**: a confident-but-wrong classification produces no observable artifact beyond the wrong posting itself, invisible to the audit trail unless a human catches it retrospectively.

This relocation framing is corroborated by independent work. Luong Tuan & Sanyal (2026) describe the equivalent phenomenon in their neurosymbolic enterprise architecture as the *inverse parametric knowledge effect*: the value of ontological grounding is highest precisely at the coverage boundary, where LLM parametric knowledge is weakest (Vietnamese financial regulation: $\Delta = +.29$; English financial domains: $\Delta = +.12$). A system well-grounded in the core of the ontology reveals its failure mode at the periphery – which is exactly the observable error behavior the CRA is designed to capture, not suppress. The neuro-symbolic financial reasoning architecture of [arXiv:2603.04663] (2026) similarly identifies that shifting from probabilistic retrieval to deterministic grounding compresses financial hallucination to 1.2%, with residual error concentrating at domain boundaries – confirming the relocation pattern as a general property of constrained-output financial systems, not an artifact of Kontablo’s specific design.

9.4 The Engineering Reframe: From Intractable to Tractable

The classical framing – “how do we stop the model from hallucinating?” – is not only intractable as stated, but directs engineering effort at the wrong target. Hallucination in

the general sense is an active area of model research with no near-term resolution that guarantees the classification accuracy a ledger requires across all inputs and distribution shifts.

The productive reframe, which Kontablo’s architecture instantiates, is:

How do we design the harness so that (a) the output space the agent can emit is bounded by the ontology, making confabulation of non-existent accounts architecturally impossible; (b) decisions within the bounded space that can be resolved by deterministic logic are resolved without inference; and (c) cases outside coverage or below confidence threshold are escalated to a human co-signer rather than silently committed?

This is a tractable engineering problem with identifiable components:

- Schema-constrained tool interfaces that enforce UUID membership
- The Three-Tier Resolution Pipeline (Section 6), which maximizes the proportion of decisions resolved deterministically before reaching the model
- The Deterministic Boundary Library (Section 6.3), which validates even Tier 3 proposals against ontological predicates before ledger write
- The Co-responsibility Architecture (Section 10), which governs the out-of-coverage escalation path

Kontablo’s architectural principle 5 – “determinism over stochasticity wherever movable” (ADR-009) – formalizes this reframe as a standing design driver: whenever a decision can be migrated from stochastic inference to deterministic logic, it must be. The harness is the system-level implementation of this principle. The ontology-as-constraint is its primary instrument.

9.5 The Co-responsibility Architecture as Residual Error Management

The Co-responsibility Architecture (ADR-008, Section 10) serves two functions (Section 12.1): it is the foundational locus of non-delegable legal accountability for every transaction, and — the aspect this subsection isolates — it is the **residual error management layer**. The discussion here concerns only the second, reactive function; it presupposes, and does not replace, the foundational accountability the human carries on every record regardless of whether it is flagged.

Where the three-tier pipeline and the Deterministic Boundary Library resolve decisions within coverage, the CRA governs what happens at and beyond the coverage boundary. When the harness cannot deterministically resolve a transaction – because no Tier 1 or Tier 2 rule matches, and Tier 3 returns confidence below threshold – the CRA routes the transaction to human review, enforces justification logging, and records the outcome in an append-only audit trail (the specification calls for cryptographic sealing of this trail; v0.x implements append-only logging).

Within this reactive scope, the human accountant is not a rubber-stamp approver of AI output. For the cases that reach review, they are the **resolver of the residual error population**: cases genuinely beyond automated coverage that require human judgment.

The CRA’s inconsistency flagging mechanism ensures that every such resolution is recorded and attributable, operationalizing the ISACA shared-responsibility model and the EU AI Act’s requirements for high-risk AI systems in financial contexts (ISACA, 2024; EU AI Act, 2025). This is precisely the human-in-the-loop anchoring that recent work on the AI assurance accountability gap identifies as the mechanism preserving professional responsibility when autonomous systems cannot bear it (Tucker & Kapoor, 2025).

This framing implies a property of the system over time that applies **only to the reactive function**: as the ontology expands – more jurisdictions mapped, more transaction types covered, more boundary predicates encoded – the population of out-of-coverage cases shrinks and the *error-resolution* load decreases. The CRA, operating on a shrinking residual, imposes progressively less reactive review burden. This does *not* mean human involvement trends to zero: the foundational accountability and context-governance function (Section 12.1) is invariant and per-transaction. What a maturing ontology changes is the *mix* of human work — less time spent resolving misclassifications, with no reduction in the legal responsibility the human carries for the record. The coverage boundary itself never reaches zero either; novel transactions, regulatory changes, and jurisdictional edge cases continually renew the residual. The design accepts both realities and manages them, rather than claiming to eliminate the human.

9.6 Future Extension: Per-Client Determinism (ADR-011)

The harness principle admits a natural extension to the per-client layer. A Client-Specific Determinism Agent (proposed in ADR-011) would observe an individual client’s transaction history, identify recurring patterns currently resolved via Tier 3 inference, and propose client-specific Tier 1 or Tier 2 rules for those patterns. This extends the harness locally, without modifying the shared ontology.

The same co-responsibility constraint applies: no client-specific rule is activated without explicit human co-signature from the client’s designated accountant. The same data sovereignty constraint applies: the agent operates entirely on client infrastructure.

ADR-011 is explicitly deferred to Phase 3+ and is documented here to establish its architectural lineage: it is an application of the same harness design principle at a finer (per-client) granularity, not a new concept. The iterative loop of ADR-009 – continuously identifying decisions that can be migrated from inference to deterministic logic – extends naturally and explicitly to the client layer once production deployments provide the empirical base required.

9.7 Neurosymbolic Coupling Maturity: Positioning Kontablo

Luong Tuan & Sanyal (2026) propose a six-level maturity model for neurosymbolic coupling in enterprise AI systems:

L0 – Ungrounded LLM operates without ontological context.

L1 – Context-Injected Ontology provides prompt context (input-side).

L2 – Discovery-Constrained Ontology filters available tools and skills.

L3 – Process-Gated Ontology enforces approval gates and escalation during execution.

L4 – Output-Validated Ontology validates and constrains LLM outputs post-generation.

L5 – Closed-Loop Full bidirectional coupling: ontology constrains inputs, guides reasoning, validates outputs, and evolves from agent experience.

Kontablo v0.1.0 operates at **L2–L3**: the Three-Tier Resolution Pipeline restricts tool discovery to valid UUID space (L2), and the Co-responsibility Architecture enforces human co-signature and audit trail for low-confidence outputs (L3). The ontology-as-constraint mechanism described in Section 9.2 – UUID membership validation before ledger write – is a hard L2 enforcement, not merely a probabilistic guardrail.

L4–L5 are explicitly deferred to future work. A full output-side validator would check that the semantic reasoning chain leading to a UUID is ontologically consistent, not merely that the UUID itself exists. This requires translating the agent’s reasoning path into a machine-checkable form against the Domain Ontology – an active research problem (Luong Tuan & Sanyal, 2026, Section 7). The architectural lineage for this extension is clear, and it is the natural next step once the L2–L3 production baseline is validated by the CPA community feedback the SSRN/Zenodo release is intended to solicit. Framing Kontablo at L2–L3 is not a concession but a precision: it allows the specific contribution of each level to be evaluated independently, and it places output-side validation in the correct temporal and engineering sequence.

10 Accounting for the Autonomous Agent Economy

10.1 The Next Infrastructure Tier: AI Agents as Entities

In the impending Machine-to-Machine (M2M) Agent Economy—driven by frameworks like Google’s *Agent2Agent* (A2A) and the *Agent Payments Protocol* (AP2)—AI agents act as autonomous economic actors. These agents negotiate, authorize, and execute high-frequency financial payments using cryptographically signed digital mandates.

11 The Role of Kontablo in the Agentic Economy

11.1 The A2A/AP2 Standard Integration

As autonomous software agents begin to negotiate, authorize, and execute transactions without human intervention, they require a subledger layer that is:

1. **Decentralized and Cryptographically Verifiable:** the specification targets ledger states validatable through cryptographic signatures and consensus rules (a design requirement; the v0.x reference implementation does not yet implement signing or consensus).
2. **Semantically Immutable:** The agent cannot redefine the meaning of financial events (e.g., classifying an expense as capital expenditure to artificially boost short-term profitability metrics).

Kontablo acts as the foundational subledger layer for the emerging **Agent Payments Protocol (AP2)** and **Agent2Agent (A2A)** networks. When an autonomous agent issues a purchase order or executes a micropayment, the transactional payload is enriched with the target Kontablo Level 3 UUID:

Payload = { Tx_ID, Amount, Currency,
Sender_Agent_ID, Recipient_Agent_ID, Kontablo_UUID }

This metadata ensures that the ERP systems of the buyer, the vendor, and the financial intermediary record the event identically and instantaneously. It eliminates the need for post-transaction bank reconciliation, as the payment message contains the exact posting coordinates for both general ledgers.

11.2 Model Context Protocol (MCP) Integration

To let AI agents interact with corporate accounting systems through a stable semantic contract rather than a bespoke database schema, the Kontablo *specification* defines a standardized **Model Context Protocol (MCP)** server interface. This is design intent, not a shipped feature: as of v0.x the reference implementation exposes its capabilities over a REST API (`api/rest`, `api/src`) and a deterministic gRPC core (`api/grpc/server.py`); the MCP server is specified but not yet implemented, and is tracked as roadmap work. Once shipped, the interface exposes tools that allow LLM agents to:

- Query the current semantic structure of the corporate chart of accounts.
- Request mapping proposals for new vendor invoices or billing items.
- Perform real-time compliance audits on draft journal entries prior to ledger write-in.

By utilizing the MCP standard, any agent equipped with a compatible client can obtain mapping proposals and validation results through structured tool calls, without learning the custom relational database schema of the underlying ERP; it interacts entirely with the stable Kontablo Graph abstraction layer. Consistent with the scope boundary of Section 12.4, Kontablo is a semantic mapping and validation layer, not a book of record: the agent does not write into a Kontablo-proprietary general ledger. The resulting transaction is recorded in the ERP through the connector — or, on the roadmap, in a thin UUID-native subledger that synchronizes to an ERP of record — with Kontablo supplying the validated posting coordinates rather than holding the authoritative ledger itself.

11.3 Cryptographic Liability and Agent_ID Tracking

In an autonomous economy, tracing liability is a critical legal requirement. When an accounting error, fraud, or transaction violation occurs, it must be possible to attribute the action to a specific autonomous agent.

The specification requires every write-in event to the Kontablo subledger to carry a payload signed by the initiating agent’s private key, with the initiating **Agent_ID** stored as tamper-evident metadata in the transactional record (the v0.x implementation records **Agent_ID** without cryptographic enforcement; signature verification is roadmap work). If the Co-responsibility Architecture flags an inconsistency, the audit trail points directly to the cryptographic identity of the responsible model. This enables forensic attribution: when an error or violation occurs, the audit trail identifies exactly which agent signed the transaction and under whose operational authority. Legal accountability is not transferred

by this attribution — the human operator who configured the system and approved its parameters retains primary legal responsibility and retroactive veto over all transactions the automated process produced. The `Agent_ID` record narrows the scope of human review; it does not replace it.

11.4 High-Frequency Micro-Transaction Aggregation

Autonomous agents often engage in high-frequency, low-value transactions (e.g., purchasing compute cycles, API calls, or sensor data). Booking each of these micro-transactions individually into a traditional ERP database would cause database lockups and excessive gas/processing costs.

Kontablo solves this through a **Micro-Transaction Aggregation** mechanism. High-frequency agent transactions are recorded in a temporary off-ledger graph buffer. The engine aggregates these edges periodically (e.g., hourly) using transitive summing:

$$\sum_{i=1}^N \text{Tx}_i(\text{Agent}_A \xrightarrow{\text{compute}} \text{Agent}_B) \xrightarrow{\text{collapse}} \text{Single Ledger Posting}$$

The collapsed sum is then written to the main General Ledger using the corresponding Level 3 UUID (`expense.admin` or `revenue.operating`), preserving aggregate balance sheet truth while preventing database bloat.

By providing a stable API ingestion layer that abstracts away the tax-code-specific labels of 195 jurisdictions, Kontablo effectively creates a “Financial Universal Service Bus” (USB) for autonomous agentic commerce. This allows agents to operate with zero friction across national borders while maintaining 100% statutory auditability for human CPAs.

12 Ontology Governance and the Co-responsibility Paradigm

12.1 The Non-Delegability Premise

The Co-responsibility Architecture rests on a premise that is legal before it is technical: **responsibility for a financial record cannot be delegated to a non-human agent, whether deterministic or stochastic.** A software artifact cannot be a defendant. It cannot be fined, sanctioned, struck off a professional register, or imprisoned. As the University of Waterloo Centre for Information System Assurance frames the accountability crisis in AI assurance, “you cannot fine an algorithm, imprison a neural network, or shame a dataset” (Tucker & Kapoor, 2025). The legal accountability for every posted transaction therefore rests, necessarily and at all times, on an identifiable human professional.

This premise has a consequence that is easy to state and important not to lose: **the human role in Kontablo is foundational, not residual.** It is present in every transaction as the locus of legal accountability, not merely invoked when the automated tiers fail. We distinguish two analytically separate human functions, because conflating them produces the false inference that the human becomes dispensable as automated coverage improves:

1. **Foundational accountability and context governance (invariant).** The human professional is the legal bearer of responsibility for the record and the

guarantor of the *integrity and continuity of context*. These are two distinct properties. *Integrity* is a point-in-time property: that the jurisdictional rules, effective-date regulations, and entity-specific facts the system reasons over are complete and correct for the transaction at hand. *Continuity* is a temporal property: that this context persists coherently across transactions and reporting periods. The distinction is not cosmetic. A stochastic language model is *stateless* — it carries no memory between invocations and therefore possesses no temporal continuity of its own. Continuity must consequently be supplied by the system rather than the model: it is materialized deterministically by the harness (append-only ontology versioning, effective-date tracking, and the audit trail described below, whose cryptographic sealing is specified as roadmap work) and held to account by the human. This function does not diminish as the ontology matures; it applies to every transaction regardless of confidence score. Crucially, accountability for context integrity and continuity is not the same as manual re-verification of every transaction: their *maintenance* is largely a systemic, deterministic responsibility of the harness, while *accountability* for them remains human. This is the shift from point-in-time audit to continuous “living compliance” (LSE Business Review, 2026).

2. **Residual error resolution (reactive).** The human resolves the cases that automated tiers cannot — the boundary population analysed in Section 9.5. This load *does* decrease as coverage expands.

The architecture below operationalizes both functions. The shrinkage argument made later in this paper (Section 9.5) applies strictly to function (2). Function (1) — legal accountability and context governance — is invariant by construction and is the reason the CRA can never reduce to zero human involvement, however complete the ontology becomes.

12.2 The Co-responsibility Architecture (CRA)

To prevent “AI-only audit loops”—where automated systems silently hide booking errors or fraud from human supervisors—Kontablo implements a **Co-responsibility Architecture (CRA)**. This architecture formalizes the governance boundary between human operators and AI models, ensuring that neither party can modify financial records without leaving a verifiable audit trail. The CRA is, in the vocabulary of contemporary **AI governance**, a domain-specific human-in-the-loop control: it instantiates the accountability, human-oversight, and record-keeping obligations that frameworks such as the NIST AI Risk Management Framework, ISO/IEC 42001, and the EU AI Act place on high-risk AI systems (ISACA, 2024; EU AI Act, 2025), specialized to the accounting ledger where the accountable party is a named professional rather than an organization. Consistent with Section 12.1, the AI is positioned throughout as a *proposer and instrumenter* — it generates informed suggestions, attaches audit and provenance metadata, and improves operational efficiency — never as the bearer of accountability for the resulting record.

The CRA operates as a multi-stage control loop:

1. **AI Proposes:** The semantic mapping engine analyzes incoming transactions and proposes a Level 3 target node ($V_{\text{universal}}$), attaching a confidence score (θ) and a natural language justification.

2. **Human Disposes:** A human accountant or domain expert reviews the proposal. The human has the authority to approve the proposed mapping or override it with a different Level 3 classification.
3. **Rule Engine Evaluates:** If the human overrides the proposal, the deterministic rule engine verifies whether the new mapping violates any core ontology constraints (e.g., nature consistency, statement boundaries, or liquidity limits). If a violation is detected, the engine does not block the transaction—which would halt business operations—but instead forces the human to enter a justification.

12.3 Database Persistence and Flagging Mechanics

The outputs of the CRA are persisted directly within the database schema of the ERP integration (e.g., as custom fields in Frappe/ERPNext’s `Journal Entry Account DocType` or SAP B1’s user-defined fields).

Two key metadata fields are appended to every mapped transaction record:

- `inconsistency_flag`: A boolean field that is set to `true` if the mapping violates a deterministic rule.
- `inconsistency_note`: A text field containing the specific rule violation description (e.g., “Nature Mismatch: local debit account mapped to credit node `equity.capital`”) concatenated with the human operator’s written justification for the override.

The specification requires that, once written, these fields be cryptographically hashed and sealed as part of the transaction block, so that any post-facto change to a mapping — for example, to cover up an irregularity — produces a hash mismatch in the audit ledger. The v0.x reference implementation provides append-only logging of these fields; cryptographic sealing (hybrid classical and post-quantum signatures over version hashes) is specified as roadmap work and is not yet implemented.

12.4 Scope and Enforcement Boundary

A recurring question determines whether the CRA is a governance guarantee or a decoration: *on which transactions does it actually run?* Kontablo is a **semantic mapping and validation layer**, not a book of record. The ERP (ERPNext, Odoo, SAP B1, NetSuite) remains the system of record; Kontablo resolves each local account to a universal node, runs the deterministic boundary checks, and writes the resulting mapping and CRA metadata back into the ERP record (the custom fields of the previous subsection) or returns them over its API. It does not maintain a shadow ledger the ERP is unaware of. Three integration modes therefore define the surface the CRA governs:

1. **Connector-mediated.** The connector subscribes to the host ERP’s write events (e.g. Frappe/ERPNext document hooks, SAP B1 UDF triggers) so that a posting — whether typed into the ERP UI, imported from a spreadsheet, or pushed by an integration — is routed through the mapping and the boundary checks at write time.

2. **API-submitted.** An autonomous agent or upstream service calls the mapping API directly (the agent-native path of Section 9); the CRA runs inline before the transaction is handed to the ERP.
3. **Batch / import.** Bulk uploads (CSV exports, migration loads) are processed through the import stream (Section 5.3) and validated as a set.

The honest boundary is this: **the CRA enforces only on transactions that reach one of these paths.** A journal entry posted directly in the ERP through a write path the connector does *not* subscribe to is governed only retroactively — at the next validation sweep or consolidation run — not at the instant of posting. Completeness of enforcement is therefore an integration property, not an automatic one: it holds exactly to the extent that the connector intercepts every write path of the host system. We state this explicitly because conflating “Kontablo validates mappings” with “Kontablo governs every transaction in the enterprise” is the ambiguity an implementer would otherwise discover late. For entries that bypass the live hooks, Kontablo’s guarantee is *detectability on sweep* — the same deterministic checks, applied in batch, flag the inconsistency and quarantine it from the consolidated statement — not prevention at keystroke.

Two authorization modes, one accountable human. The CRA does not require a human to hand-approve every line. Deterministic resolutions (Tier 1 exact-code and Tier 2 rule matches, Section 8) post under *standing ex-ante authorization*: the responsible accountant approves the deterministic policy — the code maps and the boundary rules — once, configures it, and retains veto and the right to inspect or reverse any posting it produced. These postings are logged and fully reviewable but are not individually signed, exactly as an accountant today authorizes a recurring posting rule without re-signing each instance it fires. Only the residual — flagged boundary violations, Tier 3 low-confidence proposals, and explicit human overrides — requires *per-transaction disposition*: the human selects or confirms the mapping and, where a deterministic rule is contradicted, enters the written justification persisted in **inconsistency_note**. The accountant’s “signature” is thus two records, not one: a standing authorization of the deterministic policy, and a per-exception sign-off on each flagged item. This is what is meant by a “mandatory human review *pathway* for every proposal”: the pathway and the veto exist for every transaction, while the act of review is spent where it has value — on the exceptions — rather than rubber-stamping the deterministic majority. Legal accountability remains undivided and human in both modes (Section 12.1).

12.5 Anti-Corruption and Embezzlement Mitigation

In legacy corporate environments, fraud often occurs through the collusion of internal accounting staff who “rubber-stamp” misclassified journal entries (e.g., booking personal expenditures as business-related “Administrative Expenses” or masking asset withdrawals as “Intercompany Receivables”).

By inserting the CRA between the general ledger database and the user interface, Kontablo establishes an automated, independent “witness.” Because the deterministic rules cannot be bribed or pressured, any fraudulent mapping override will automatically flag itself in the audit database. When external auditors perform their review, they do not need to sample thousands of transactions randomly; they simply query for all records where `inconsistency_flag == true`. This structural transparency dramatically

increases the difficulty of internal embezzlement and enforces compliance at the database write-in layer.

13 Validation: Mass Consolidation Simulation

To test the ontology’s robustness against the three crises identified in Section 1.1, we developed a mass-consolidation engine and exercised it on **synthetic** trial balances — 10 distinct entities across 9 countries representing diverse regulatory environments. The data is simulated, not drawn from production ledgers; what the simulation validates is the mapping and consolidation *architecture*, not field performance.

13.1 Consolidation Simulation Engine

The simulation models a multinational holding company consolidating trial balances from 10 distinct operational entities across 9 jurisdictions: Mexico (MX, two entities), Brazil (BR), France (FR), Panama (PA), Ecuador (EC), Colombia (CO), Vietnam (VN), Nigeria (NG), and Saudi Arabia (SA). Each subsidiary ledger is a synthetic trial balance written in the trial-balance export format of one of four source ERPs (ERPNext, Zoho Books, Odoo, SAP Business One) — the files mirror each ERP’s export layout but were never exported from a live ERP instance — carrying localized charts of accounts (real statutory codes from SAT, SPED-style, PCG, PUC, VAS, SOCPA numbering), varying currencies, and diverging tax treatment. Every entity’s trial balance balances exactly (total debits equal total credits in local currency), and the consolidated ledger satisfies the accounting identity $A = L + E + (R - X)$ with escalated lines carried in an explicit suspense bucket. The full artifact — fixtures, parsers, resolution pipeline, and results — is committed and regenerable with a single deterministic command (`scripts/consolidation_v1_initial_run.py`; artifacts in `research/experiments/consolidation_v1/`); no API keys or model calls are required. This initial run was subsequently extended into a larger, fully reproducible validation matrix (75 entities, 68 jurisdictions; Section 13.5).

13.2 The 10-Point Complexity Scale

To measure the difficulty of mapping local charts of accounts to the universal Level 3 ontology, we established a **10-point Complexity Scale**. The scale rates jurisdictions based on four primary vectors: (1) alignment with IFRS, (2) strictness of mandatory national charts of accounts, (3) complexity of electronic invoicing and transactional metadata requirements, and (4) currency volatility/inflation treatment.

Kontablo’s full ontology covers all **195 UN-recognized sovereign states**; the IFRS Foundation’s own jurisdiction survey covers 169 jurisdictions (IFRS Foundation, 2025). Table 1 presents a **stratified representative sample of 21 jurisdictions** drawn from that full mapping, selected to illustrate the range of complexity tiers from IFRS-pure low-friction environments to IAS 29 hyperinflationary stress cases. Within this sample, the average complexity score is **5.6 / 10** (arithmetic mean over the 21 rows; see table), reflecting the deliberate stratification across all difficulty tiers rather than a random draw.

Jurisdiction	Complexity (1–10)	Primary Complexity Driver	Validation Status
<i>Tier 1–2: IFRS-pure / minimal local divergence</i>			
United Kingdom (UK)	2/10	Full IFRS adoption; common-law flexibility	Fully Mapped
Canada (CA)	2/10	IFRS aligned with CPA Canada guidance	Fully Mapped
Australia (AU)	2/10	AASB equivalent to IFRS verbatim	Fully Mapped
<i>Tier 3–4: Moderate divergence</i>			
United States (US)	3/10	US-GAAP (FASB ASC); convergence ongoing	Fully Mapped
India (IN)	4/10	Ind-AS (IFRS converged, carve-outs apply)	Fully Mapped
Germany (DE)	4/10	HGB commercial code statutory divergence	Mapped (IKR Overlay)
UAE (AE)	4/10	Federal corporate tax introduced 2023 (MoF)	Fully Mapped
<i>Tier 5: Mandatory national chart + digital-tax metadata</i>			
Japan (JP)	5/10	J-GAAP (ASBJ); mandatory corporate-tax schedules	Fully Mapped
Mexico (MX)	5/10	Strict SAT 5-digit catalog tagging	Fully Mapped
Colombia (CO)	5/10	Rigid PUC numeric hierarchy	Fully Mapped
Panama (PA)	5/10	DGI/SMV dual-standard reporting	Fully Mapped
South Africa (ZA)	5/10	SARS guidelines; IFRS adopted via SAICA	Fully Mapped
<i>Tier 6: Codified national plans + inflation</i>			
France (FR)	6/10	Mandatory Plan Comptable Général (PCG)	Mapped (PCG Overlay)
Spain (ES)	6/10	PGC 8-group structure; strict numbering	Fully Mapped
China (CN)	6/10	CAS standard divergence; VAT reform	Fully Mapped
Russia (RU)	6/10	RAS standard; transitioning to IFRS	Mapped (RAS Overlay)
Turkey (TR)	6/10	TAS inflation accounting rules (IAS 29 border)	Fully Mapped
<i>Tier 7: Complex tax regimes / strict government charts</i>			
Brazil (BR)	7/10	SPED e-reporting; ICMS/PIS/COFINS tax credits	Fully Mapped
Saudi Arabia (SA)	7/10	Zakat religious tax base (SOCPA)	Mapped (Zakat Overlay)
Vietnam (VN)	7/10	Strict VAS 2014 government chart mandates	Mapped (VAS Overlay)
<i>Tier 10: IAS 29 hyperinflation stress case</i>			
Venezuela (VE)	10/10	IAS 29 hyperinflation; dual-currency parallel rates	Stress Tested

Table 1: Stratified representative sample (21 of 195 mapped jurisdictions) illustrating the full complexity-tier range of the Kontablo ontology. Complexity scores reflect: IFRS alignment, mandatory national chart strictness, digital-tax metadata requirements, and currency/inflation treatment. Mean complexity for this sample: 5.6/10. US-GAAP source: FASB ASC (fasb.org); Ind-AS source: MCA notification 2015 (mca.gov.in); J-GAAP source: ASBJ (asb.or.jp). Full 195-jurisdiction dataset available in the project repository.

13.3 Coverage and Data Quality

The AI-augmented graph successfully mapped 100% of the sample transaction sets to Kontablo’s core ‘asset.current.cash’ and ‘asset.current.receivables’ universal IDs. By dynamically applying local-to-target foreign exchange conversions and semantic rate overrides, the engine demonstrated that Kontablo can seamlessly aggregate localized financial ledgers into a unified global balance sheet. A separate, reproducible transaction-frequency benchmark (`scripts/coverage_benchmark.py`, dataset and provenance in `research/coverage_benchmark/`) estimates that the 30 Level 3 accounts cover **~94% of routine transaction volume by posting-line count** (and ~87% counted by whole transaction). We are explicit about the evidentiary status of this figure: it is a *model-based estimate* computed over a transparently-parameterized synthetic frequency distribution, not a measurement of a proprietary ledger corpus. The distribution’s structure is grounded in the standard double-entry “special journals” taxonomy (Weygandt, Kieso & Kimmel; Horngren et al.), under which the bulk of postings flow through four high-volume journals (Sales, Cash Receipts, Purchases, Cash Payments) whose transaction types map entirely onto core nodes, while the General Journal carries only infrequent adjusting entries. The committed artifact replaces an earlier, unbacked appeal to a large but uncommitted set of ledger exports; substituting a licensed real-ledger corpus is a documented next step that leaves the script unchanged. The substantive finding is robust to the exact weights: extreme ledger granularity in national standards (often hundreds of accounts) is mostly reserved for rare statistical or edge-case events.

13.4 Coverage Expansion and the Principled Residual

The ~94% figure is a coverage-by-volume measure of a *minimum core* of 30 nodes, not a ceiling. Because transaction frequency follows a long-tail (Pareto-type) distribution — a small number of account types carry the bulk of posting volume, while a long tail of specialized types occurs rarely — coverage is raised not by redesigning the architecture but by **extending the ontology with additional Level 3 nodes for the next tier of transaction types**. The same benchmark identifies these next-tier types directly from the residual: in descending volume order they are payroll net-pay and statutory-withholding liabilities, contract/deferred revenue (IFRS 15), tax withheld at source on income, and non-trade receivables. Adding exactly these **four** nodes — a 34-account *extended core*, kept separate from the frozen 30-account minimum core — lifts modeled volume coverage from ~94% to **~99%** (~97% by whole transaction). This both confirms and sharpens the earlier rough estimate (which guessed “a few dozen” nodes): in this profile the bulk of the residual collapses onto a handful of high-frequency payroll- and tax-related liabilities. The final percentage points correspond to genuinely idiosyncratic, jurisdiction-specific, or novel instruments, which we deliberately leave to escalation rather than core growth.

We deliberately do **not** claim, or target, 100% coverage. Two reasons, one principled and one architectural. *Principled*: a universal standard that asserted total coverage would be epistemically overconfident — new instruments (tokenised assets, CBDC settlement, novel Islamic-finance structures) and new jurisdictional rules continuously regenerate the tail, so a closed claim of completeness would be false the moment it was made. A residual uncertainty band is the honest representation of an open, evolving domain. *Architectural*: Kontablo already has a designated mechanism for the residual — the semantic coverage boundary (Section 9.3) and the Co-responsibility Architecture (Section 9), which route out-of-coverage transactions to human review rather than forcing a low-confidence posting.

Coverage expansion shrinks the residual; the CRA guarantees that whatever residual remains is escalated and made observable, not silently mismapped. The objective is therefore not coverage = 100% but coverage $\rightarrow$ high with a *bounded, instrumented, human-resolved* remainder.

This raises a governance question the benchmark answers directly: *under what criteria may the core grow, and when must it not?* We codify the answer as an explicit admission policy (ADR-015). The ontology is structured in three layers — a **frozen** 30-node minimum core (the citable Pareto headline), a **bounded** extended core that grows only under the policy, and an open-ended residual handled by overlays and CRA escalation that is *never* forced into the core. A transaction type earns a core node only if it simultaneously (i) captures material, *measured* posting volume (the benchmark is the gate, with a working threshold near 0.5% of total volume), (ii) recurs across many jurisdictions under a clean IFRS anchor (otherwise it is a local overlay, not a universal node), (iii) admits an unambiguous deterministic boundary — a fixed nature, statement class and IFRS tag with no collision — (iv) is not already decomposable into existing nodes, and (v) is a mature, stable instrument. Because posting volume is Pareto-distributed, this gate is **self-limiting**: after the extended core, almost every further candidate falls below the volume threshold, so the core converges on a small number of dozens of nodes rather than the hundreds found in national charts. Growth by intuition or by request is thereby replaced with growth by measured evidence, and the standing tension between a small, elegant core and broad coverage is resolved without compromising either.

13.5 Expanded Multi-Jurisdiction Validation

What “195 jurisdictions” means. Coverage is recorded in an explicit manifest of all 195 sovereign states — the 193 UN member states plus the two UN observer states (Holy See, State of Palestine) — in `core/schemas/jurisdiction_coverage.yaml1`. The claim is precise and two-layered. *Every* one of the 195 is covered by Kontablo’s universal IFRS-anchored layer, because each Level-3 node carries an `ifrs_tag`; this is the mapping basis for the **135** jurisdictions that do not mandate a national numeric chart of accounts (typically common-law / IFRS-direct regimes — including IFRS-converged systems such as Japan, South Korea, Indonesia, Thailand and Switzerland, and IFRS-adopting CIS states such as Armenia, Georgia, Azerbaijan, Kyrgyzstan and Tajikistan, that prescribe reporting rules but no legally mandated numeric chart — where an entity designs its own chart and there is no statutory code to map). The remaining **60** jurisdictions impose a statutory chart and additionally receive a local-code overlay; for **56** of these we populate and exercise primary-source-cited Tier-1 code sets, including the SYSCOHADA chart shared identically by all 17 OHADA member states (one verified chart, 17 jurisdictions); the Spanish PGC (RD 1514/2007), the Portuguese SNC (Portaria 218/2015), the Peruvian PCGE, the Chinese CAS (MOF) and the Belgian PCMN (CNC/CBN); the Moroccan CGNC, the Algerian SCF and the Romanian chart (OMFP 1802/2014); the Czech (Vyhláška 500/2002), Slovak, Hungarian (Act C/2000) and Bulgarian national charts; the Ukrainian (Order 291), Kazakh, Tunisian (PCT 1997) and Belarusian charts; the Serbian, Croatian (RRiF), Slovenian and Austrian (EKR) charts; the Luxembourg PCN and the French PCG (also applied in Monaco); the Moldovan and Greek (EGLS/ELP) charts; and single-country charts such as Panama’s (11 cited codes) and Mexico’s SAT. The four statutory-chart jurisdictions not yet carrying Tier-1 codes (Poland, Uzbekistan, Turkmenistan, Laos) mandate a chart whose account numbers we could not confirm from

a primary source — Poland’s are not legally fixed — and are left Tier-1-pending rather than transcribed speculatively (boundary B1).

Boundary conditions of the coverage claim. To keep the 195 figure honest rather than rhetorical, the manifest is governed by an explicit construction protocol, recorded in its metadata: (B1) no fabricated codes — a jurisdiction is marked Tier-1 only when a primary-source-cited code set exists, and a deterministic filter (a real statutory code must contain at least one digit) excludes descriptive text placeholders such as “Cash” or “Vorsteuer” from the Tier-1 index (14 such placeholders were demoted, and the affected common-law jurisdictions fall back to the IFRS-tag basis); (B2) two distinct claims are reported separately — `mapping_mode=statutory_chart` (a mandated chart *exists*) versus `tier1_codes_available` (we hold *verified* codes), so “statutory without codes” never masquerades as Tier-1 coverage; (B3) IFRS-direct is a genuine coverage mode, not a gap, since the IFRS tag is the complete mapping basis where no local code exists; (B4) multi-state multipliers are used only where a chart is genuinely shared (SYSCOHADA across 17 states; the EU, having no shared chart, is covered per country); (B5) jurisdictions that apply IFRS without a mandated chart (e.g. Chile, Uruguay, Costa Rica) are left IFRS-direct rather than asserting a chart that does not exist; (B6) IFRS adoption status derives from the IFRS Foundation profiles, with unprofiled states marked as such. The coverage claim is thus the conjunction of a universal IFRS-anchored layer and a verifiable, conservatively-bounded statutory overlay — not an assertion of 195 transcribed national charts.

To test the pipeline beyond the initial proof, we ran a broader, fully reproducible consolidation (script `scripts/mass_consolidation_v2.py`; artifacts in `research/experiments/consolidation`). The trial balances are *generated directly from the ontology and chart families themselves*: for every jurisdiction with real `local_codes` (inline or via a chart family), an entity is constructed from those exact codes, so Tier-1 exact lookup is exercised against genuine mapping data rather than hand-authored examples. The run processes **366 local account entries from 75 entities across 68 jurisdictions**, spanning all five statement classes (assets, liabilities, equity, revenue, expense) and fourteen languages across five writing systems (Latin: Spanish, Portuguese, French, Italian, Polish, Indonesian, Vietnamese, German, Turkish, English; Greek; Cyrillic: Russian; Arabic; Hangul: Korean), and populates **25 of the 30 universal Level-3 nodes**. To keep the run deterministic and reproducible, the semantic AI fallback (Tier 3) is *not* exercised; instead, any entry the deterministic tiers cannot resolve is counted as **escalated** (routed to human review by the CRA). The run is therefore a test of pipeline correctness and boundary behaviour on a controlled dataset, not an out-of-sample accuracy benchmark.

Three results are worth emphasising. First, the **four escalated entries are all genuine coverage-boundary cases**, not failures: three are v0.3-roadmap instrument types not yet ratified in the operational ontology (cryptocurrency holdings, carbon credits, and a Zakat provision — all carried in the YAML `pending_accounts` block without a confirmed `nature`), and one is a jurisdiction-specific German tax-driven reserve (*Sonderposten mit Rücklageanteil*) with no clean IFRS counterpart. The engine *escalated rather than mismapped* each of them — the empirical instantiation of the principled residual of Section 13.4. Second, the Co-responsibility Architecture was tested against a deliberately injected error catalog spanning **five deterministic accounting invariants** — debit/credit nature, the liquidity boundary (liquid asset to non-current), statement class (a P&L name forced onto a balance-sheet node and vice versa), VAT direction

Metric	Result	
Sovereign states in coverage manifest	195	
universal IFRS-anchor layer	195	
statutory-chart overlay	60	
IFRS-direct (mapping via IFRS tag)	135	
Tier-1 code sets populated (cited)	56 / 60 statutory	
Entities consolidated	75	
Jurisdictions exercised in run	68	
IAS 29 hyperinflation dual-rate cases	4 (VE, AR, LB, TR)	
Local account entries processed	366	
Tier 1 (exact local-code lookup)	310	(84.7%)
Tier 2 (deterministic keyword rules)	46	(12.6%)
Deterministic coverage (Tier 1+2)	97.3%	
Escalated to human via CRA	4	(1.1%)
Injected-error catalog detected & quarantined by CRA	8 / 8	
distinct invariant classes exercised	5	
(nature, liquidity, statement, VAT direction, equity/liability)	12 flags	
Descriptive placeholders excluded from Tier-1 (B1)	14	
Ontology code-collisions: detected → corrected → remaining	4 → 0	
Distinct Kontablo nodes populated	25 / 30	
Consolidated total assets (FX-normalized)	USD 14,746,038	

Table 2: 195-jurisdiction coverage manifest and the deterministic validation run (75 entities, 68 jurisdictions), with Tier-1 trial balances generated from the ontology’s own `local_codes` and statutory chart families (e.g., SYSCOHADA for the 17 OHADA states). Coverage is measured by entry count; FX normalization to USD uses simulated 2026 rates, with parallel-rate overrides for hyperinflationary entities. The hyperinflation cases exercise IAS 29’s dual-rate (official vs. parallel) translation divergence only; CPI restatement of non-monetary items, which IAS 29 requires before translation, is not exercised by this harness.

(input forced onto the output node and vice versa), and equity/liability classification. It **detected and quarantined all eight injected errors** (twelve flags in total), holding them for human review rather than posting them. No injected error passed silently, and none reached the consolidated ledger. The scope of this result must be stated precisely: the nature and liquidity checks (invariants one and two) run unconditionally on every resolved mapping, but the statement-class, VAT-direction, and equity/liability checks run only on entries pre-labeled as upstream proposals (`forced_id`), and each injected name was written to contain vocabulary the corresponding check matches. The catalog therefore demonstrates that each invariant class is *mechanically wired end-to-end* — a flagged entry is provably quarantined before posting — not a detection-rate measurement; recall on plausible-but-wrong proposals whose names evade the keyword vocabulary is unmeasured and is future work. Third, and most consequential for the development of the standard, the run **closed a corrective loop on the ontology itself**. It surfaced four latent local-code collisions — a single jurisdiction code mapped to two different Kontablo nodes (MX 509 → interest/fx-loss; RU 68.2 → input/output VAT; VN 311 → short-/long-term debt; VN 635 → interest/fx-loss). Each was then corrected against the primary national standard (SAT *código agrupador* Anexo 24 for MX; the Russian chart Order 94n, which separates input VAT in account 19 from VAT settlements in 68-02; and Vietnam’s Circular 200/2014, which assigns long-term borrowings to account 341). The VN 635 case was recognised not as a typo but as a genuine *granularity mismatch* — VAS aggregates interest and FX losses under one statutory code — and resolved by mapping the code to its dominant leaf and documenting the aggregation rather than fabricating a spurious distinct code. After correction, re-validation reported **zero** collisions. This is a concrete instance of Kontablo’s central thesis operating on Kontablo: a machine-verifiable standard is a self-checking one. The harness functions not only as a mapping engine but as a continuous data-quality check on the standard it implements; the full find→fix→re-validate record is archived with the run artifacts.

Limitation: what the 97.3% figure does and does not measure. The deterministic-coverage figure is an *internal-consistency and data-quality result*, not a field-accuracy measurement. The Tier-1 entries are generated by iterating over the ontology’s own committed `local_codes` and chart families and are then resolved against an index built from the same data, so Tier-1 success is guaranteed by construction except where the data itself is defective — which is precisely what makes the run a self-check: the four code collisions it surfaced are real defects it caught in its own answer key. The Tier-2 entries are author-written names matched against author-written multilingual rules, and the four escalations are deliberately out-of-vocabulary frontier instruments. None of this exercises the noise of real ledger exports — abbreviations (“A/R Trade”), department-suffixed codes (“1100-01-DEPT7”), or misspellings — and Tier-2 substring matching should be expected to degrade on such input. A perturbation experiment (typos, prefixed codes, abbreviations) over the committed fixtures is planned future work; until it is committed, the honest reading of 97.3% is “the pipeline mechanics and the ontology’s code data are mutually consistent,” not “97% of real-world ledger lines resolve deterministically.”

13.6 Non-Anglophone Jurisdictions and the Value of Deterministic Grounding

The initial run exposes a pattern across the non-Anglophone jurisdictions rated high on the complexity rubric of Table 1 (Vietnam 7/10, Saudi Arabia 7/10) that aligns with the **Inverse Parametric Knowledge Effect (IPKE)** identified in independent empirical work (Luong Tuan & Sanyal, 2026; Section 8.5): ontological grounding adds most value precisely where LLM pretraining coverage of the target domain is structurally thinnest.

The committed run (`research/experiments/consolidation_v1/results.json`) shows the two sides of this effect. For Saudi Arabia (SOCPA with Zakat overlay), the deterministic tiers resolve only 5 of 12 lines — one ontology-carried SOCPA code at Tier 1, four Arabic names matched by the multilingual Tier 2 rules — and the remaining 7 lines escalate to human review, the highest escalation rate in the run. (In these reproducible runs, lines the deterministic tiers cannot resolve are *escalated* rather than passed to the Tier 3 semantic fallback, which would require model calls; in a deployment, these escalations are exactly the lines that Tier 3 would receive.) For Vietnam, by contrast, the ontology carries the VAS Circular 200 statutory code set, and 12 of 13 lines resolve at Tier 1 — committed deterministic grounding substitutes entirely for the parametric knowledge the model lacks. The contrast is the point: where neither committed codes nor keyword rules cover a jurisdiction, the residual routes to semantic inference or to a human, and without the Deterministic Boundary Library’s hard rejection of UUID non-members those would be the lines most likely to produce silent misclassifications. The deterministic constraint layer is therefore not merely a quality enhancement; it is an operational requirement for jurisdictions where the model’s parametric knowledge is sparse.

This observation generalises. Of the 195 jurisdictions mapped, those with non-Latin charts and thin English-language regulatory documentation (SYSCOHADA francophone Africa, Korean K-GAAP, Ukrainian P(S)BO, Arabic SOCPA, Vietnamese VAS) are precisely the jurisdictions where Kontablo’s constraint layer provides the highest marginal safety guarantee over an unconstrained RAG or pure-LLM approach. The universality of coverage, argued structurally in Appendix D, is reinforced empirically by the finding that the hardest jurisdictions to map are also the ones where hard constraints matter most.

14 Labor Market Implications: The Accountant Re-composition Hypothesis

A complete account of Kontablo’s implications cannot ignore its effects on the human practitioners it interacts with. This section addresses a question that the architecture raises but does not answer by design: does a system that automates transaction classification and batch reconciliation reduce the demand for accountants? The answer, grounded in labor-market evidence, is more precise than a binary yes or no.

14.1 The Intra-Occupation Split

The most commonly cited concern — that AI threatens the accounting profession — conflates two occupations that the U.S. Bureau of Labor Statistics (BLS) tracks and projects separately. The BLS Occupational Outlook Handbook (2025 edition) projects

accountants and auditors at +5% employment growth from 2024 to 2034, with 124,200 annual job openings and a 2.0% unemployment rate — strong by any labour-market standard. It projects *bookkeeping, accounting, and auditing clerks* at −6% over the same period, explicitly attributing the decline to automation of routine data-entry and reconciliation tasks (BLS, 2025a; BLS, 2025b).

The World Economic Forum Future of Jobs Report (2025) corroborates the split at a global scale: accounting and bookkeeping clerks rank seventh among the fastest-declining occupations worldwide, with an estimated decline of approximately 20% by 2030 (WEF, 2025). The same report categorises accountants and auditors — the judgment-intensive tier — as stable or growing.

This intra-occupation split is the empirically grounded starting point for understanding Kontablo’s labor-market position. Kontablo’s Tier 1 and Tier 2 automation directly targets the task set that drives clerk-level decline: exact-code lookup, regex-based disambiguation, and batch reconciliation. These are the tasks the BLS and WEF identify as automation-susceptible. The CPA-level tasks — cross-jurisdictional interpretation, regulatory judgment, AI output oversight, and advisory work — are precisely those that Tier 3 cannot resolve deterministically and the Co-responsibility Architecture escalates to human review.

14.2 The Co-responsibility Architecture as a Structural Demand Signal

The Co-responsibility Architecture (CRA, Section 10) carries a labor-market implication that is architectural rather than incidental. The CRA is not an optional human-review layer added for optics; it is the mechanism that manages the residual error population (Section 9.5) — the transactions that are genuinely beyond automated resolution. Every escalation through the CRA invokes a human accountant who must exercise substantive judgment: interpreting a flagged inconsistency in the context of local regulatory rules, overriding a low-confidence Tier 3 mapping with documented justification, or resolving a jurisdictional edge case for which no Tier 1 or Tier 2 rule yet exists.

This is the architectural mechanism underlying what we term the **Accountant Recomposition Hypothesis**: Kontablo does not reduce the accounting profession — it recomposes it. The system automates what is automatable (Tier 1/2 execution tasks) and structurally requires human expertise for what is not (the CRA residual). The minimum competency for the role that remains is higher, not lower.

This claim is consistent with field evidence. The Inverse Parametric Knowledge Effect (Section 8.5) implies that the hardest jurisdictions — VAS-Vietnam, SOCPA-Saudi Arabia, SYSCOHADA-francophone Africa, IAS 29 hyperinflationary contexts — generate the highest Tier 3 escalation rates and therefore the greatest demand for accountants with genuine multi-jurisdictional expertise. The universality of Kontablo’s 195-jurisdiction coverage does not flatten this demand; it exposes and concentrates it at the boundary.

14.3 Empirical Support: What the Market is Signalling

Several independent data sources are consistent with the recomposition hypothesis as of mid-2026.

Skill premium and complementarity. The PwC Global AI Jobs Barometer (2025), an analysis of close to a billion job advertisements across six continents, found a 56%

wage premium for roles requiring AI skills, and 7.5% year-on-year growth in postings requiring AI skills across sectors (PwC reports the growth economy-wide, not for finance specifically). Mäkelä & Stephany (2024) analysed 12 million job vacancies and found that AI-exposed roles are twice as likely to require complementary cognitive and interpersonal skills, with a measurable spillover effect elevating skill requirements even in roles not directly AI-augmented (arXiv:2412.19754). These findings support the interpretation that automation at one layer of the task stack raises the value of the cognitive skills above it — the pattern Kontablo’s architecture instantiates by design.

Acute CPA shortage. Robert Half’s 2026 Accounting and Finance Hiring Guide reports that CPA-level roles take an average of 73 days to fill, 41% longer than in prior survey periods, with respondents describing the market for experienced accountants as “acutely constrained” (Robert Half, 2026). The AICPA Trends in the Supply of Accounting Graduates (2025) confirms strong hiring intent across firms of all sizes, with the pipeline challenge described as structural (reduced CPA-exam sitting rates) rather than AI-driven (AICPA, 2025).

Practitioner-level evidence. ICAEW’s 2026 mid-tier firms study (*Evolution of Mid-Tier Accountancy Firms*), which surveyed leaders from 35 mid-tier UK firms in February–March 2026, found that 68% of respondents believed AI would reduce demand for some early-career accounting roles, but 83% stated it would not result in fewer roles overall, and 71% reported that AI had enabled their teams to move up the value chain toward advisory and strategic work (ICAEW, 2026). The sample is small and bounded to the mid-tier segment, so this is directional practitioner sentiment rather than population-level evidence; its pattern — automation compressing entry-level execution while expanding senior-level advisory demand — is nonetheless the recomposition dynamic at the firm level.

14.4 The Honest Tension: Long-Run Residual Shrinkage

Epistemic integrity requires naming a structural tension in the recomposition hypothesis. The *reactive* function of the CRA operates on the population of cases that automated tiers cannot resolve, and that function’s load decreases as deterministic coverage expands (Section 9.5). If one mistakenly equated the human role with this reactive function alone, the conclusion would follow that a fully converged ontology requires less high-level human review — the opposite of the recomposition hypothesis.

This apparent tension dissolves once the two human functions are kept distinct (Section 12.1). What shrinks with coverage is the *error-resolution* load, not the human role. The foundational function — non-delegable legal accountability for every record and governance of context integrity — is invariant and per-transaction. A non-human agent cannot bear legal responsibility for a financial record (Tucker & Kapoor, 2025; Meditari Accountancy Research, 2024), so even a hypothetical fully-converged ontology still requires a qualified human accountable for each posting. The recomposition is therefore in the *nature* of the work: automation removes the execution and error-resolution tasks, leaving the accountability, judgment, and context-governance tasks that demand senior competency.

The empirical counter to this is that the coverage boundary is not static. New regulatory regimes emerge — Turkey adopted IAS 29 hyperinflation accounting in 2022; Lebanon in 2020; Argentina’s restatements under *NIC 29* continue to evolve (PwC, 2026; EY, 2025). New agentic economy instruments (tokenised securities, CBDC settlement,

multi-agent payment rails) generate transaction types for which no Tier 1 or Tier 2 rule yet exists (Davidovic & Tourpe, 2026; AP2, 2026). The acute CPA shortage (Robert Half, 2026) is direct evidence that the existing human capacity for judgment-intensive accounting work is already insufficient to meet demand — *before* the Kontablo-level abstraction layer expands the frontier of addressable jurisdictions and instruments.

The recomposition hypothesis is therefore not falsified by the long-run shrinkage argument: the residual renews structurally. The claim is not that Kontablo is a net job creator, but that it is a task-level recomposer that raises the floor of required competency and concentrates demand at the expertise tier the labour market is already unable to supply.

14.5 Implications for Professional Development

The recomposition hypothesis, if correct, has a normative implication for accounting education and certification. Competency frameworks that train practitioners primarily for Tier 1/2 execution tasks — routine classification, reconciliation, journal entry processing — will prepare graduates for the portion of the profession that Kontablo and equivalent systems automate. Competency frameworks that train for multi-jurisdictional interpretation, AI-output oversight, and boundary-case resolution will prepare graduates for the residual that the CRA requires.

This is consistent with decisions already visible in the market: the ACCA has updated its syllabus to include AI ethics and sustainability as core competencies from 2024 (ACCA, 2024); ICAEW’s 2026 survey shows 71% of firms moving practitioners toward advisory work. The architectural design of Kontablo’s CRA makes the same implication concrete at the system level: the accountant who interacts with the system is not reviewing routine postings but resolving the cases that are genuinely hard — and that role requires the expertise tier that is currently in shortage.

The design intent of the CRA is not to create employment. It is to ensure that every automated decision with a non-trivial confidence distribution is resolved by a qualified human, with durable, logged attribution. The labor-market consequence — that this structural requirement sustains demand for high-level practitioners — is a property of the architecture, not an aspirational claim.

15 Conclusion and Future Work

In this paper, we presented Kontablo, an open-source, IFRS-anchored universal accounting ontology designed to resolve the semantic fragmentation of localized ledger systems (the “Accounting Babel” problem). By representing the financial universe as a directed multi-dimensional property graph rather than a rigid tree hierarchy, Kontablo provides a stable, jurisdiction-agnostic subledger for both multinational ERP consolidation and the high-frequency Machine-to-Machine (M2M) Agentic Economy.

To enable secure, automated translation, we introduced a three-tier mapping resolution strategy (exact lookup, disambiguation rules, and semantic AI fallback) and bound it to a Co-responsibility Architecture (CRA). The CRA ensures that human override authority is preserved while automatically auditing deviations through append-only database inconsistency flags (cryptographic sealing of the trail is specified as roadmap work). We validated this framework on synthetic, ERP-formatted trial balances using a reproducible mass-consolidation engine spanning 75 entities across 68 jurisdictions

(extended from an initial nine-country proof), demonstrating that a minimum core of 30 Level 3 accounts can represent an estimated $\sim 94\%$ of routine transaction volume (rising to $\sim 99\%$ with a 34-account extended core, per the reproducible benchmark in `research/coverage_benchmark/`) while handling IAS 29 dual-rate hyperinflation translation (index restatement of non-monetary items is not exercised) and local overlays (such as VAT and Zakat). The same run also functioned as a corrective loop on the standard, surfacing and resolving four latent local-code collisions in the ontology, and its Co-responsibility Architecture quarantined every injected inconsistency across five deterministic accounting invariants.

Future research and development will focus on the following roadmap areas:

1. **Global Expert Validation Phase:** Conducting structured validation interviews with CPAs and accounting standard-setters across OECD and non-OECD nations to refine the Level 3 core.
2. **Production ERP Connectors:** Building production-grade, two-way sync connectors for NetSuite and SAP S/4HANA using the Tree-to-Graph linearization protocol.
3. **AP2 and A2A Protocol Standards:** Formulating the official RFC specification for embedding Kontablo Level 3 UUIDs within the transactional payloads of the Agent Payments Protocol (AP2).
4. **Specialized Industry Overlays:** Expanding the ontology with specialized Level 3/4 overlays for banking (IFRS 9), insurance (IFRS 17), and agriculture (IAS 41).
5. **High-Frequency Agentic Tracking:** Standardizing `Agent_ID` fields and micro-transaction batching mechanisms within the Kontablo Graph.

The Kontablo Universal Accounting Ontology is published under the Business Source License 1.1 (BSL 1.1), converting to Apache 2.0 after four years; connectors to open-source ERP platforms are licensed Apache 2.0. The full research data, mass-consolidation engine, and AI mapping services are available for public auditing and collaborative extension.

Acknowledgements

The author thanks the open-source communities behind ERPNext/Frappe, Odoo, and the IFRS Foundation’s public taxonomy releases, whose freely accessible standards made the 23-jurisdiction empirical analysis tractable. The jurisdictional coverage of Vietnamese (VAS) and Saudi Arabian (SOCPA) standards benefited from AI-assisted bilingual semantic mapping; all resulting mappings were reviewed by the author.

This work was conducted independently. No external funding was received. The Kontablo Project is an initiative of Praxia, Cuenca, Ecuador, currently in incorporation.

16 Appendix: Level 3 Account Catalog

The following table lists a comprehensive selection of the core Level 3 accounts defined in the Kontablo v0.2.0 schema, illustrating the mapping of localized codes from five

representative jurisdictions (Mexico SAT, Colombia PUC, Panama DGI, Vietnam VAS, France PCG).

Level 3 ID	UUID	Nature	IFRS Tag	Local Code Mappings (MX, CO, PA, VN, FR)
asset.current.cash	...0101	Debit	CashAndCashEquivalents	MX: 101, CO: 1105, PA: 1.1.01, VN: 111, FR: 512
asset.current.bank	...0102	Debit	CashAndCashEquivalents	MX: 102, CO: 1110, PA: 1.1.02, VN: 112, FR: 512
asset.current.receivables	...0104	Debit	TradeAndOtherReceivables	MX: 105, CO: 1305, PA: 1.1.05, VN: 131, FR: 411
asset.current.vat_input	...0001	Debit	CurrentTaxAssetsCurrent	MX: IVA Acred., CO: 2408, PA: 1.1.12, VN: 1331, FR: 445660
asset.current.inventory	...0106	Debit	CurrentInventories	MX: 151, CO: 1430, PA: 1.1.08, VN: 152, FR: 371
asset.current.prepaid	...0001	Debit	CurrentPrepayments	MX: 122, CO: 1705, PA: 1.1.10, VN: 121, FR: 488
asset.noncurrent.ppe	...0110	Debit	PropertyPlantAndEquipment	MX: 181, CO: 1516, PA: 2.1.01, VN: 211, FR: 211
asset.noncurrent.rou_assets	...0001	Debit	RightOfUseAssets	MX: N/A, CO: N/A, PA: Lease Asset, VN: N/A, FR: 215
asset.noncurrent.intangibles	...0002	Debit	IntangibleAssetsOtherThanGW	MX: 191, CO: 1630, PA: 2.2.01, VN: 213, FR: 2800
asset.noncurrent.goodwill	...0003	Debit	Goodwill	MX: N/A, CO: 1605, PA: Goodwill, VN: N/A, FR: 207
liability.current.payables	...0201	Credit	TradeAndOtherPayables	MX: 201, CO: 2205, PA: 3.1.01, VN: 331, FR: 401
liability.current.vat_output	...0203	Credit	CurrentTaxLiabilities	MX: IVA Trasl., CO: 2408, PA: N/A, VN: 3331, FR: 445710
liability.current.tax	...0205	Credit	CurrentTaxLiabilities	MX: 2030, CO: 2404, PA: 3.1.05, VN: 3334, FR: 444
liability.current.accrued	...0207	Credit	OtherCurrentProvisions	MX: 270, CO: 2805, PA: 3.1.08, VN: 333, FR: 3900
liability.noncurrent.debt	...0211	Credit	OtherNonCurrentFinLiabilities	MX: 311, CO: 2105, PA: 4.1.01, VN: 311, FR: 164
equity.capital	...0008	Credit	IssuedCapital	MX: 300, CO: 3105, PA: 5.1.01, VN: 411, FR: 101
equity.retained	...0009	Credit	RetainedEarnings	MX: 305, CO: 3705, PA: 5.1.05, VN: 421, FR: 110
revenue.operating	...0301	Credit	Revenue	MX: 401, CO: 4135, PA: 6.1.01, VN: 511, FR: 701
expense.cogs	...0404	Debit	CostOfSales	MX: 501, CO: 6105, PA: 7.1.01, VN: 632, FR: 607
expense.admin	...0501	Debit	AdministrativeExpenses	MX: 504, CO: 5205, PA: 7.2.01, VN: 642, FR: 627
expense.depreciation	...0503	Debit	DepreciationExpense	MX: 506, CO: 5160, PA: 7.2.05, VN: 214, FR: 681
expense.interest	...0505	Debit	FinanceCosts	MX: 509, CO: 5305, PA: 7.3.01, VN: 635, FR: 661
expense.tax	...0507	Debit	TaxExpense	MX: 510, CO: 5405, PA: 7.4.01, VN: 821, FR: 695

Table 3: Comprehensive list of Level 3 Accounts showing mapped local codes across representative jurisdictions.

Appendix B: Jurisdiction Data Availability

The research analyzed 195 jurisdictions across four economic clusters. Primary machine-readable data sources include the **Mexican SAT Chart of Accounts (2023)**, the **French PCG (2022)**, and the **Brazilian Plano Referencial (2024)**. Vietnamese (VAS) and Saudi Arabian (SOCPA) data were analyzed using bilingual AI-assisted semantic mapping to reconcile non-Latin nomenclatures.

Appendix C: Tree-to-Graph-to-Tree — Linearization Protocol

The following figure illustrates the full round-trip of the Kontablo linearization protocol: a local ERP chart of accounts (flat tree) is imported into the Kontablo Universal Graph (multi-dimensional), where each account participates simultaneously in multiple independent dimensions. For any reporting or ERP export target, the graph is linearized by selecting a single dimension as the primary axis, recovering a flat output tree (e.g., an IFRS Balance Sheet). No information is lost in either direction: the graph retains the full multi-dimensional classification while the tree provides downward compatibility with legacy ERP schemas.

Appendix D: IFRS-Pure Jurisdictions and the Universality Thesis

A notable finding from the 195-jurisdiction mapping effort is that the final three jurisdictions to reach complete coverage — Estonia (EE), Botswana (BW), and Namibia (NA) — share a common characteristic: all three are **pure IFRS adopters with no mandatory national chart of accounts**.

This is not incidental. These jurisdictions require no adaptation layer whatsoever: entities operating under full IFRS adoption map directly onto the Kontablo Level 3 minimum-core without any jurisdiction-specific override. The ontology covers them through its IFRS-derived core, not through a localization file with custom codes.

The pattern generalizes. The 195 mapped jurisdictions fall into two structural categories:

1. **Mandatory-chart jurisdictions:** A national authority mandates a specific chart of accounts (e.g., France PCG, Mexico SAT, Brazil SPED, the 15 OHADA SYSCO-HADA members). These require a localization mapping that cross-references national codes to Kontablo UUIDs. The mapping is deterministic and machine-verifiable.
2. **IFRS-pure jurisdictions:** No mandatory chart exists; entities self-select account structures within IFRS constraints. These map directly to the universal core with no intermediary. Coverage is structurally guaranteed, not empirically discovered.

The implication for the universality thesis (Section 2) is direct: *no sovereign accounting jurisdiction escapes the ontology*. Either the jurisdiction has a mandatory chart (in which case a deterministic mapping exists) or it does not (in which case the IFRS-derived core applies directly). The two cases are exhaustive and mutually exclusive. A hypothetical new sovereign state adopting IFRS from inception would be covered on day zero, without any new localization work.

This finding was discovered empirically during the coverage expansion: the three “last” countries to be added were precisely those that required the least jurisdictional work, confirming the theoretical boundary of the problem.

Special accounting contexts — IAS 29 hyperinflation (Venezuela, Lebanon, Argentina, Turkey, Sudan, South Sudan, Iran, Malawi, and Haiti as at June 2026; PricewaterhouseCoopers, 2026), Islamic finance instruments (Saudi Arabia, Qatar, Kuwait, Bahrain,

Pakistan, Brunei, Sudan), distribution-only corporate income tax (Estonia, Latvia, Georgia), and no-VAT jurisdictions (Hong Kong, Brunei, Vanuatu, Antigua & Barbuda, and others) — required dedicated account extensions within the universal graph, demonstrating that contextual completeness, not structural novelty, is the remaining engineering challenge at the boundary of global coverage.

Selected Bibliography

- IFRS Foundation. (2024). *IFRS Taxonomy 2024*.
- IFRS Foundation. (2025). *Who Uses IFRS Accounting Standards?* <https://www.ifrs.org/use-around-the-world/use-of-ifrs-standards-by-jurisdiction/> [Profiles for 169 jurisdictions.]
- International Accounting Standards Board. *IAS 21 — The Effects of Changes in Foreign Exchange Rates*. IFRS Foundation. [Foreign-currency transactions recorded at the spot rate at the transaction date; average-rate practical expedient permitted where rates do not fluctuate significantly.]
- European Securities and Markets Authority. (2023). *European Single Electronic Format (ESEF) Reporting Manual*. ESMA. [Mandated inline-XBRL (iXBRL) disclosure format for EU-regulated issuers, tagged against the IFRS Taxonomy.]
- SAP SE. (2024). *SAP S/4HANA Finance: the Universal Journal (ACDOCA)*. [Single unified line-item table consolidating GL, controlling, and reporting for multi-GAAP, multi-currency postings within one instance.]
- Oracle Corporation. (2024). *Oracle ERP: Primary, Secondary, and Reporting Ledgers*. [Parallel primary (e.g. US-GAAP) and secondary (e.g. IFRS) ledgers within a single instance with automated conversion journals.]
- ISACA. (2024). *AI Shared Responsibility Model: A Governance Framework*.
- EU AI Act. (2025). *Governance of High-Risk AI Systems in Financial Markets*.
- Deloitte. (2025). *Benchmarks for Global Financial Operations: The Cost of Inefficiency*.
- Gartner. (2024). *Transforming Finance through Autonomous Systems*.
- Fintech Global. (2025). *The Inefficiency Crisis: Addressing Reconciliation Error Rates in B2B Payments*.
- Davidovic, S., & Tourpe, H. (2026). *How Agentic AI Will Reshape Payments*. IMF Notes, 2026/004. <https://www.imf.org/en/publications/imf-notes/issues/2026/04/22/how-agentic-ai-will-reshape-payments-575560>
- Google Cloud / Linux Foundation. (2025). *Agent2Agent (A2A) Protocol*. Donated to Linux Foundation June 2025. <https://a2a-protocol.org/latest/>
- Google / FIDO Alliance. (2025). *Agent Payments Protocol (AP2)*. Donated to FIDO Alliance May 2026. <https://ap2-protocol.org/>

- World Economic Forum. (2025). *Agentic Commerce and the Future of M2M Transactions*.
- OECD. (2023). *Tax Administration 2.0: Structural Digitalization*.
- Luciani, C. (2026). *The Accounting Babel: Cross-Border Semantic Fragility*.
- G7 Cyber Expert Group. (2026, January). *Roadmap for Coordinating the Transition to Post-Quantum Cryptography in the Financial Sector*. Co-chaired by U.S. Treasury and Bank of England. <https://home.treasury.gov/news/press-releases/sb0355>
- National Institute of Standards and Technology. (2024, August). *NIST Releases First 3 Finalized Post-Quantum Encryption Standards* (FIPS 203 ML-KEM, FIPS 204 ML-DSA, FIPS 205 SLH-DSA). <https://www.nist.gov/news-events/news/2024/08/nist-releases-first-3-finalized-post-quantum-encryption-standards>
- Agand, P. (2026). *Neuro-Symbolic Financial Reasoning via Deterministic Fact Ledgers and Adversarial Low-Latency Hallucination Detector*. arXiv:2603.04663. <https://arxiv.org/abs/2603.04663> [Universal Fact Ledger; deterministic grounding compresses financial hallucination to 1.2%.]
- Thanh Luong Tuan, & Abhijit Sanyal. (2026). *Ontology-Constrained Neural Reasoning in Enterprise Agentic Systems: A Neurosymbolic Architecture for Domain-Grounded AI Agents*. arXiv:2604.00555v5. <https://arxiv.org/abs/2604.00555> [Authors listed in full per arXiv; Vietnamese family/given split not asserted. 1,800-run experiment; Vietnamese domains show 2× ontology lift; L0–L5 coupling maturity model; Foundation AgenticOS platform, 22 industry verticals.]
- *PHANTOM: A Benchmark for Hallucination Detection in Financial Long-Context QA* (2025). *NeurIPS 2025 Datasets and Benchmarks Track*. <https://openreview.net/forum?id=5YQAoOS3Hm> [Authorship to be confirmed at camera-ready.]
- Sharma, K., Kumar, P., & Li, Y. (2025). OG-RAG: Ontology-Grounded Retrieval-Augmented Generation for Large Language Models. *EMNLP 2025*. [+55% fact recall, +40% response correctness via ontology-anchored hypergraph retrieval.]
- Willard, B. T., & Louf, R. (2023). Efficient Guided Generation for Large Language Models. arXiv:2307.09702. [Grammar-constrained decoding; syntactic output bounds without fluency loss.]
- Luo, L., et al. (2025). Graph-Constrained Reasoning: Faithful Reasoning on Knowledge Graphs with Large Language Models. *ICML 2025*. arXiv:2410.13080. [KG-Trie constrains LLM decoding to valid KG paths at the token level.]
- Garcez, A. d., & Lamb, L. C. (2023). Neurosymbolic AI: The 3rd Wave. *Artificial Intelligence Review* 56: 12387–12406. DOI:10.1007/s10462-023-10448-w.
- PricewaterhouseCoopers. (2026, June). *Hyperinflationary Economies as at June 2026*. https://viewpoint.pwc.com/gx/en/pwc/in-briefs/2026/ib_int202610.html

- Ernst & Young. (2025, November). *IFRS Developments Issue 242: Hyperinflationary Economies*. <https://www.ey.com/content/dam/ey-unified-site/ey-com/en-gl/technical/ifrs-technical-resources/documents/ey-gl-developments-242-hyperinflationary-economies.pdf>
- Spivak, D. I. (2012). Functorial Data Migration. *Information and Computation*. arXiv:1009.1166.
- Spivak, D. I., & Kent, R. E. (2012). Ologs: A Categorical Framework for Knowledge Representation. *PLoS ONE* 7(1): e24274.
- Ellerman, D. (2014). On Double-Entry Bookkeeping: The Mathematical Treatment. *Accounting Education* 23(5): 483–501. arXiv:1407.1898.
- Robinson, M. (2017). Sheaves are the Canonical Data Structure for Sensor Integration. *Information Fusion* 36: 208–224.
- Hansen, J., & Ghrist, R. (2019). Toward a Spectral Theory of Cellular Sheaves. *Journal of Applied and Computational Topology* 3. arXiv:1808.01513.
- Curry, J. (2014). *Sheaves, Cosheaves and Applications*. PhD thesis. arXiv:1303.3255.
- Thiéblin, E., Haemmerlé, O., Hernandez, N., & Trojahn, C. (2020). Survey on Complex Ontology Matching. *Semantic Web* 11(4). DOI:10.3233/SW-190366.
- Vauquois, B. (1968). A Survey of Formal Grammars and Algorithms for Recognition and Transformation in Machine Translation. *IFIP Congress*.
- Hasić, D., & Tannier, E. (2019). Gene Tree Species Tree Reconciliation. arXiv:1703.08950.
- Mémoli, F. (2011). Gromov–Wasserstein Distances and the Metric Approach to Object Matching. *Foundations of Computational Mathematics* 11(4): 417–487.
- Xu, H., Luo, D., Zha, H., & Carin, L. (2019). Gromov–Wasserstein Learning for Graph Matching and Node Embedding. *ICML*. arXiv:1901.06003.
- Cover, T. M. (1965). Geometrical and Statistical Properties of Systems of Linear Inequalities with Applications in Pattern Recognition. *IEEE Trans. Electron. Comput.* EC-14(3): 326–334.
- Bureau of Labor Statistics. (2025a). *Occupational Outlook Handbook: Accountants and Auditors*. U.S. Department of Labor. <https://www.bls.gov/ooh/business-and-financial/accountants-and-auditors.htm> [Projected employment change 2024–2034: +5%; 124,200 annual openings; 2.0% unemployment rate.]
- Bureau of Labor Statistics. (2025b). *Occupational Outlook Handbook: Bookkeeping, Accounting, and Auditing Clerks*. U.S. Department of Labor. <https://www.bls.gov/ooh/office-and-administrative-support/bookkeeping-accounting-and-auditing-clerks.htm> [Projected employment change 2024–2034: –6%; automation cited as primary cause.]
- World Economic Forum. (2025). *Future of Jobs Report 2025*. Geneva: WEF. [Accounting and bookkeeping clerks ranked 7th fastest-declining occupation globally; estimated –20% by 2030.]

- Institute of Chartered Accountants in England and Wales. (2026). *Evolution of Mid-Tier Accountancy Firms*. ICAEW. [Survey of leaders from 35 mid-tier UK firms, February–March 2026: 68% expect some early-career role reduction; 83% expect no net role reduction; 71% report teams moving toward advisory work.]
- Mäkelä, M., & Stephany, F. (2024). AI Complementarity and Labour Market Recomposition: Evidence from 12 Million Job Vacancies. arXiv:2412.19754. Oxford Internet Institute. [AI-exposed roles 2× more likely to require complementary cognitive skills; spillover effect across firm-level job postings.]
- PricewaterhouseCoopers. (2025). *Global AI Jobs Barometer*. PwC. [Analysis of close to a billion job ads across six continents: 56% wage premium for roles requiring AI skills; 7.5% year-on-year growth in AI-skill postings across sectors.]
- Robert Half. (2026). *2026 Accounting and Finance Hiring Guide*. Robert Half International. [CPA-level roles average 73 days to fill, 41% longer than prior survey period; market described as “acutely constrained”.]
- American Institute of CPAs. (2025). *Trends in the Supply of Accounting Graduates and the Demand for Public Accounting Recruits*. AICPA. [Strong hiring intent across firm sizes; pipeline challenge attributed to reduced CPA-exam sitting rates, not AI displacement.]
- Association of Chartered Certified Accountants. (2024). *Strategic Professional Syllabus Update: AI Ethics and Sustainability as Core Competencies, effective 2024*. ACCA. <https://www.accaglobal.com/>
- Tucker, J., & Kapoor, A. (2025, December). *The Accountability Crisis of AI in Assurance: Rethinking Professional Codes of Ethics*. University of Waterloo Centre for Information System Assurance. <https://uwaterloo.ca/uwaterloo-centre-for-information-system-assurance-accountability-crisis-ai-assurance-rethinking-professional> [“You cannot fine an algorithm, imprison a neural network, or shame a dataset”; human-in-the-loop as the anchor of professional responsibility.]
- Meditari Accountancy Research. (2024). Artificial Intelligence Legal Personality and Accountability: Auditors’ Accounts of Capabilities and Challenges for Instrument Boundary. *Meditari Accountancy Research* 32(7): 120–148. Emerald Publishing. [Auditor judgment as a non-delegable core function; against premature AI legal personhood that would dilute human accountability.]
- London School of Economics Business Review. (2026, February). *AI Governance Must Move from “Point-in-Time” Audits to “Living” Compliance*. LSE. <https://blogs.lse.ac.uk/businessreview/2026/02/19/ai-governance-must-move-from-point-in-time-audits-to-living-compliance/> [Context integrity as a continuous, not snapshot, governance responsibility.]
- Harvard Journal of Law & Technology. (2025). *Redefining the Standard of Human Oversight for AI Negligence*. <https://jolt.law.harvard.edu/digest/redefining-the-standard-of-human-oversight-for-ai-negligence>
- Gruber, T. R. (1993). A Translation Approach to Portable Ontology Specifications. *Knowledge Acquisition* 5(2): 199–220. [Canonical definition of “ontology” in the

information-science sense: a formal, explicit specification of a shared conceptualization.]

- Business Research Industry. (2026, March). *Global Financial Close Software Market Size/Share Worth USD 18.5 Billion by 2034 at an 8.1% CAGR*. [Financial close software market valued at USD 8.8 billion in 2025.] <https://www.globenewswire.com/news-release/2026/03/18/3258374/0/en/>
- Market Research (financial consolidation segment). (2025). *Financial Consolidation Software Market Size and Forecast*. [Segment valued at approximately USD 3.2 billion in 2025, CAGR $\approx 9\%$.] Figures vary by analyst house; reported here as an order-of-magnitude proxy, not a precise total.
- Gartner. (2021). *How to Improve Your Data Quality / Data Quality Market Survey*. [Poor data quality costs the average organization an estimated USD 12.9 million per year.] <https://www.gartner.com/en/data-analytics/topics/data-quality>
- Reta Vortaro (ReVo). *Esperanto dictionary: konto, kontado, kalkuli*. <https://vortaro.net> [Source for the Esperanto etymology of “Kontablo”.]




Figure 8: The Deterministic Boundary Library as a validation gauntlet. An AI-proposed mapping must satisfy a sequence of deterministic accounting invariants — the five exercised in the validation harness (nature, liquidity, statement class, VAT direction, and equity/liability classification). Any single failure quarantines the proposal for human review with a penalised confidence score; it is *not* posted to the ledger. Only a clean pass on every invariant reaches the ledger write path.




Figure 9: Locus of error before and after ontology-as-constraint is applied. Classical hallucination (arbitrary invalid outputs) is eliminated at the harness layer. Residual errors relocate to the semantic coverage boundary: unmapped transaction types, incomplete jurisdictional rules, and semantically ambiguous cases within valid UUID space. Cases outside coverage are escalated to human review rather than silently committed.




Figure 10: The Co-responsibility Architecture (CRA) feedback loop. An AI-proposed mapping is reviewed by a human accountant. If the human override violates deterministic ontology rules (e.g., changing cash nature), the system forces entry of an explanation, flags the transaction, and writes it to the immutable ledger for future auditor review.




Figure 11: Tree-to-Graph-to-Tree: the Kontablo linearization protocol, read left to right. A local ERP chart of accounts (left, orange) is a rigid 2-D tree in which each account has a single parent. On import (Tree-to-Graph), it enters the Kontablo Universal Graph (centre, blue), where each account node carries a UUID and participates simultaneously in multiple independent graph dimensions — IFRS/Balance Sheet (solid blue), Tax/Regulatory (dashed red), and Functional/Cost-Centre (dotted violet); the centre is drawn in isometric pseudo-3-D to make this multi-dimensionality explicit. For reporting or ERP export, the graph is linearized (Graph-to-Tree) by selecting a single dimension as the primary axis, producing a clean 2-D IFRS output tree (right, green) without loss of the multi-dimensional data held in the graph. To make the transformation legible, one example account (*Cash*) is traced in gold across all three stages: its local leaf **101 Cash**, its UUID node in the graph, and its IFRS output *Cash & Equivalents*; the remaining accounts are mapped in faint grey.